



**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
(ARTIFICIAL INTELLIGENCE & MACHINE LEARNING)**

B.E / SEM VII / REV 2019 'C SCHEME' / CSE-(AI&ML)
Academic Year: 2023-24

NAME	PRATHAMESH S. CHIKANKAR
BRANCH	CSE-(AI&ML)
ROLL NO.	
SUBJECT	DEEP LEARNING LAB
COURSE CODE	CSL701
PRACTICAL NO.	
DOP	
DOS	



Experiment No. 01

Aim :- To explore python libraries for deep learning e.g. Theano, TensorFlow etc.

Theory :-

1. **TensorFlow** : TensorFlow is a popular open-source library for deep learning developed by Google. It provides a flexible ecosystem for building and deploying machine learning models. TensorFlow uses a computational graph approach, where you define a graph of operations and execute them within a TensorFlow session. It offers a high-level API called Keras that simplifies the process of building and training deep learning models.

TensorFlow supports distributed computing, allowing you to train models on multiple machines or GPUs simultaneously.

TensorFlow offers hardware acceleration through its integration with specialized hardware like GPUs and TPUs (Tensor Processing Units).

- **Eager Execution:** TensorFlow 2.x introduced Eager Execution, which allows immediate evaluation of operations, making TensorFlow behave more like standard Python code.
- **tf.data API:** TensorFlow provides the tf.data API, which enables efficient and flexible data input pipelines for large datasets.
- **SavedModel:** TensorFlow's SavedModel format allows you to save the entire model, including the architecture, weights, and training configuration.
- **Custom Training Loops:** TensorFlow enables you to build custom training loops, providing greater control over the training process.

```
✓ [16] import tensorflow as tf
0s

# Define the input placeholders
a = tf.keras.Input(shape=(2,), dtype=tf.float32)
b = tf.keras.Input(shape=(2,), dtype=tf.float32)

# Define the matrix multiplication operation
c = tf.matmul(a, tf.transpose(b))

# Create a TensorFlow model
model = tf.keras.Model(inputs=[a, b], outputs=c)

# Perform matrix multiplication
x = tf.constant([[1, 2], [3, 4]], dtype=tf.float32)
y = tf.constant([[5, 6], [7, 8]], dtype=tf.float32)
result = model.predict([x, y])

print(result)

1/1 [=====] - 0s 105ms/step
[[17. 23.]
 [39. 53.]]
```



2. Keras (built on top of TensorFlow) : Keras is a high-level neural networks API that provides a user-friendly interface for building and training deep learning models.

Keras supports multiple backends, including TensorFlow, Theano, and CNTK. In this example, we'll use TensorFlow as the backend.

Essentials/features -

- >High-Level API
- >Model Subclassing
- >Callbacks
- >Pre-Trained Models
- >Transfer Learning

```
✓ 5s  import numpy as np
    from keras.models import Sequential
    from keras.layers import Dense

    # Define the training data
    X_train = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])
    y_train = np.array([0, 1, 1, 0])

    # Build the model
    model = Sequential()
    model.add(Dense(4, input_dim=2, activation='relu'))
    model.add(Dense(1, activation='sigmoid'))

    # Compile the model
    model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

    # Train the model
    model.fit(X_train, y_train, epochs=100, batch_size=1)

    # Get the final predictions
    predictions = model.predict(X_train)
    predictions = np.round(predictions)

    print("Final Predictions:")
    for i in range(len(predictions)):
        print(f"Input: {X_train[i]}, Predicted: {predictions[i][0]}")
```



```
Epoch 1/100
4/4 [=====] - 2s 6ms/step - loss: 0.6858 - accuracy: 0.2500
Epoch 2/100
4/4 [=====] - 0s 6ms/step - loss: 0.6838 - accuracy: 0.2500
Epoch 3/100
4/4 [=====] - 0s 6ms/step - loss: 0.6827 - accuracy: 0.2500
Epoch 4/100
4/4 [=====] - 0s 7ms/step - loss: 0.6814 - accuracy: 0.2500
Epoch 5/100
4/4 [=====] - 0s 8ms/step - loss: 0.6809 - accuracy: 0.2500
Epoch 6/100
4/4 [=====] - 0s 7ms/step - loss: 0.6798 - accuracy: 0.2500
Epoch 7/100
4/4 [=====] - 0s 5ms/step - loss: 0.6795 - accuracy: 0.2500
Epoch 8/100
4/4 [=====] - 0s 17ms/step - loss: 0.6785 - accuracy: 0.2500
Epoch 9/100
4/4 [=====] - 0s 10ms/step - loss: 0.6777 - accuracy: 0.2500
Epoch 10/100
4/4 [=====] - 0s 10ms/step - loss: 0.6777 - accuracy: 0.2500
✓ 5s completed at 14:34

Epoch 98/100
4/4 [=====] - 0s 4ms/step - loss: 0.6072 - accuracy: 1.0000
Epoch 99/100
4/4 [=====] - 0s 7ms/step - loss: 0.6068 - accuracy: 1.0000
Epoch 100/100
4/4 [=====] - 0s 5ms/step - loss: 0.6058 - accuracy: 1.0000
1/1 [=====] - 0s 104ms/step

Final Predictions:
Input: [0 0], Predicted: 0.0
Input: [0 1], Predicted: 1.0
Input: [1 0], Predicted: 1.0
Input: [1 1], Predicted: 0.0
```



3. Theano : Theano is a popular library for numerical computation and building deep learning models. It allows efficient computation on both CPUs and GPUs.

Theano provides a symbolic math library that enables you to define, optimize, and evaluate mathematical expressions involving multi-dimensional arrays efficiently.

Theano operates on the concept of symbolic computation. It builds a computational graph that represents the mathematical expressions as a series of operations.

Theano was designed with deep learning in mind and has been used as a foundational library by other frameworks such as Keras.

Theano's development has been discontinued, and its successor, JAX, provides similar functionality with additional features.

- **Symbolic Computation:** One of the key features of Theano is its symbolic computation approach.
- **Automatic Differentiation:** Theano provides automatic differentiation, which is essential for training neural networks through backpropagation.
- **GPU Acceleration:** Theano has built-in support for GPU acceleration, allowing you to perform computations on compatible NVIDIA GPUs.
- **Shared Variables:** Theano introduces shared variables, which allow you to allocate memory on the GPU or CPU and share it between different functions.
- **Convolution and Pooling Operations:** Theano provides built-in functions for common operations used in deep learning, including convolution and pooling operations, which are fundamental in convolutional neural networks (CNNs).
- **Speed and Efficiency:** Theano was designed to provide efficient computation, making it a suitable choice for complex mathematical operations often involved in deep learning models.
- **Deprecated Status:** As of September 2020, Theano is no longer actively developed and maintained by its original developers.

```
[ ] import theano
      import theano.tensor as T

      # Define the symbolic variables
      a = T.matrix('a')
      b = T.matrix('b')

      # Define the symbolic expression
      c = T.dot(a, b)

      # Compile the function
      matmul = theano.function([a, b], c)

      # Perform matrix multiplication
      x = [[1, 2], [3, 4]]
      y = [[5, 6], [7, 8]]
      result = matmul(x, y)

      print(result)
```

Conclusion :- We had successfully explored python libraries for deep learning such as Keras, TensorFlow, Theano and PyTorch.



3. PyTorch : PyTorch is an open-source deep learning library developed by Facebook's AI Research lab (FAIR). It provides a flexible and efficient framework for building and training neural networks. PyTorch is widely used by researchers and practitioners for various machine learning tasks due to its ease of use, dynamic computation graph, and excellent support for automatic differentiation, which is essential for training deep learning models.

Key Concepts in PyTorch :

- **Tensors:** Tensors are multi-dimensional arrays similar to NumPy arrays but with additional capabilities for GPU acceleration..
- **Autograd:** PyTorch's automatic differentiation engine. It automatically tracks and computes gradients for tensors.
- **Neural Network Module (nn.Module):** PyTorch provides a base class, nn.Module, that allows you to define neural network architectures as a series of layers or functions.
- **Loss Functions:** PyTorch includes a variety of loss functions for different tasks, such as classification (cross-entropy loss) and regression (mean squared error).
- **Optimizers:** PyTorch offers various optimization algorithms like stochastic gradient descent (SGD), Adam, RMSprop, etc., to update the model parameters during training.

```
[11] import torch
      import torch.nn as nn
      import torch.optim as optim
      from sklearn.datasets import load_iris
      from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import OneHotEncoder

      iris = load_iris()
      X, y = iris.data, iris.target
      encoder = OneHotEncoder()
      y_onehot = encoder.fit_transform(y.reshape(-1, 1)).toarray()
      X_tensor = torch.tensor(X, dtype=torch.float32)
      y_tensor = torch.tensor(y, dtype=torch.long)

      class SimpleNN(nn.Module):
          def __init__(self):
              super(SimpleNN, self).__init__()
              self.fc1 = nn.Linear(4, 10)
              self.fc2 = nn.Linear(10, 3)

          def forward(self, x):
              x = torch.relu(self.fc1(x))
              x = torch.softmax(self.fc2(x), dim=1)
              return x

      model = SimpleNN()
      criterion = nn.CrossEntropyLoss()
      optimizer = optim.Adam(model.parameters(), lr=0.001)

      for _ in range(50):
          optimizer.zero_grad()
          outputs = model(X_tensor)
          loss = criterion(outputs, y_tensor)
          loss.backward()
          optimizer.step()

          accuracy = (model(X_tensor).argmax(1) == y_tensor).float().mean().item()
          print("Accuracy:", accuracy)
```

Accuracy: 0.6600000262260437



**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
(ARTIFICIAL INTELLIGENCE & MACHINE LEARNING)**

B.E / SEM VII / REV 2019 'C SCHEME' / CSE-(AI&ML)
Academic Year: 2023-24

NAME	PRATHAMESH S. CHIKANKAR
BRANCH	CSE-(AI&ML)
ROLL NO.	
SUBJECT	DEEP LEARNING LAB
COURSE CODE	CSL701
PRACTICAL NO.	
DOP	
DOS	



Program - Output :

```
In [2]: import torch
import torch.nn as nn
import torch.optim as optim

# XOR gate inputs and outputs
X = torch.tensor([[0, 0], [0, 1], [1, 0], [1, 1]], dtype=torch.float32)
y = torch.tensor([[0], [1], [1], [0]], dtype=torch.float32)

class MLP(nn.Module):
    def __init__(self):
        super(MLP, self).__init__()
        self.fc1 = nn.Linear(2, 4) # 2 input features, 4 hidden units in the first layer
        self.fc2 = nn.Linear(4, 4) # 4 hidden units, 4 hidden units in the second layer
        self.fc3 = nn.Linear(4, 1) # 4 hidden units, 1 output

    def forward(self, x):
        x = torch.sigmoid(self.fc1(x))
        x = torch.sigmoid(self.fc2(x))
        x = torch.sigmoid(self.fc3(x))
        return x

model = MLP()
criterion = nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=0.1)

# Training the model
num_epochs = 10000
for epoch in range(num_epochs):
    optimizer.zero_grad()
    outputs = model(X)
    loss = criterion(outputs, y)
    loss.backward()
    optimizer.step()

# Evaluating the model
with torch.no_grad():
    predicted = model(X)
    predicted = torch.round(predicted) # Round predictions to 0 or 1
    print("Predicted:")
    print(predicted)

Predicted:
tensor([[0.],
       [1.],
       [0.],
       [1.]])
```



**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
(ARTIFICIAL INTELLIGENCE & MACHINE LEARNING)**

B.E / SEM VII / REV 2019 'C SCHEME' / CSE-(AI&ML)
Academic Year: 2023-24

NAME	PRATHAMESH S. CHIKANKAR
BRANCH	CSE-(AI&ML)
ROLL NO.	
SUBJECT	DEEP LEARNING LAB
COURSE CODE	CSL701
PRACTICAL NO.	
DOP	
DOS	



Output:

CO DLL_EXP3.ipynb ☆

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

[1] import numpy as np

{x} ✓ 0s

▶ # Simulated data for linear regression
np.random.seed(42)
X = np.random.rand(50, 1)
y = 2 * X + 1 + 0.1 * np.random.randn(50, 1)

Helper function to compute gradient for linear regression
def compute_gradient(xi, yi, W, b):
 error = (W * xi + b - yi)
 gradient_W = 2 * xi * error
 gradient_b = 2 * error
 return gradient_W, gradient_b

▼ Stochastic Gradient Descent

{x} ✓ 0s

▶ # Optimization algorithms
def sgd(X, y, W, b, learning_rate, num_epochs):
 for _ in range(num_epochs):
 for xi, yi in zip(X, y):
 gradient_W, gradient_b = compute_gradient(xi, yi, W, b)
 W -= learning_rate * gradient_W
 b -= learning_rate * gradient_b
 return W, b

Hyperparameters
learning_rate = 0.01
batch_size = 10
num_epochs = 50
beta = 0.9

Initial parameters
W_initial, b_initial = np.random.randn(), np.random.randn()

Apply optimization algorithms
W_sgd, b_sgd = sgd(X,y,W_initial,b_initial,learning_rate,num_epochs)
print("SGD - Final Parameters: W =", W_sgd, "b =", b_sgd)

□ SGD - Final Parameters: W = [0.10320444] b = [0.98500344]



Mini-Batch Gradient Descent

{x}

```
[4] def mini_batch_gd(X, y, W, b, learning_rate, batch_size, num_epochs):
    num_samples = len(X)
    for _ in range(num_epochs):
        for i in range(0, num_samples, batch_size):
            X_batch, y_batch = X[i:i+batch_size], y[i:i+batch_size]
            gradient_W_avg = np.mean([compute_gradient(xi, yi, W, b)[0] for xi, yi in zip(X_batch, y_batch)])
            gradient_b_avg = np.mean([compute_gradient(xi, yi, W, b)[1] for xi, yi in zip(X_batch, y_batch)])
            W -= learning_rate * gradient_W_avg
            b -= learning_rate * gradient_b_avg
    return W, b

# Hyperparameters
learning_rate = 0.01
batch_size = 10
num_epochs = 50
beta = 0.9

# Initial parameters
W_initial, b_initial = np.random.randn(), np.random.randn()

# Apply optimization algorithms
W_mini_batch, b_mini_batch = mini_batch_gd(X,y,W_initial,b_initial, learning_rate,batch_size,num_epochs)
print("Mini Batch GD - Final Parameters: W =", W_mini_batch, "b =", b_mini_batch)
```

Mini Batch GD - Final Parameters: W = 0.7911081086681657 b = 1.5708653450938141

Momentum Gradient Descent

```
[5] def momentum_gd(X, y, W, b, learning_rate, beta, num_epochs):
    velocity_W, velocity_b = 0, 0
    for _ in range(num_epochs):
        for xi, yi in zip(X, y):
            gradient_W, gradient_b = compute_gradient(xi, yi, W, b)
            velocity_W = beta * velocity_W + (1 - beta) * gradient_W
            velocity_b = beta * velocity_b + (1 - beta) * gradient_b
            W -= learning_rate * velocity_W
            b -= learning_rate * velocity_b
    return W, b

# Hyperparameters
learning_rate = 0.01
batch_size = 10
num_epochs = 50
beta = 0.9

# Initial parameters
W_initial, b_initial = np.random.randn(), np.random.randn()

# Apply optimization algorithms
W_momentum, b_momentum = momentum_gd(X,y,W_initial,b_initial, learning_rate,beta,num_epochs)
print("Momentum GD - Final Parameters: W =", W_momentum, "b =", b_momentum)
```

Momentum GD - Final Parameters: W = [1.94629625] b = [1.02554776]

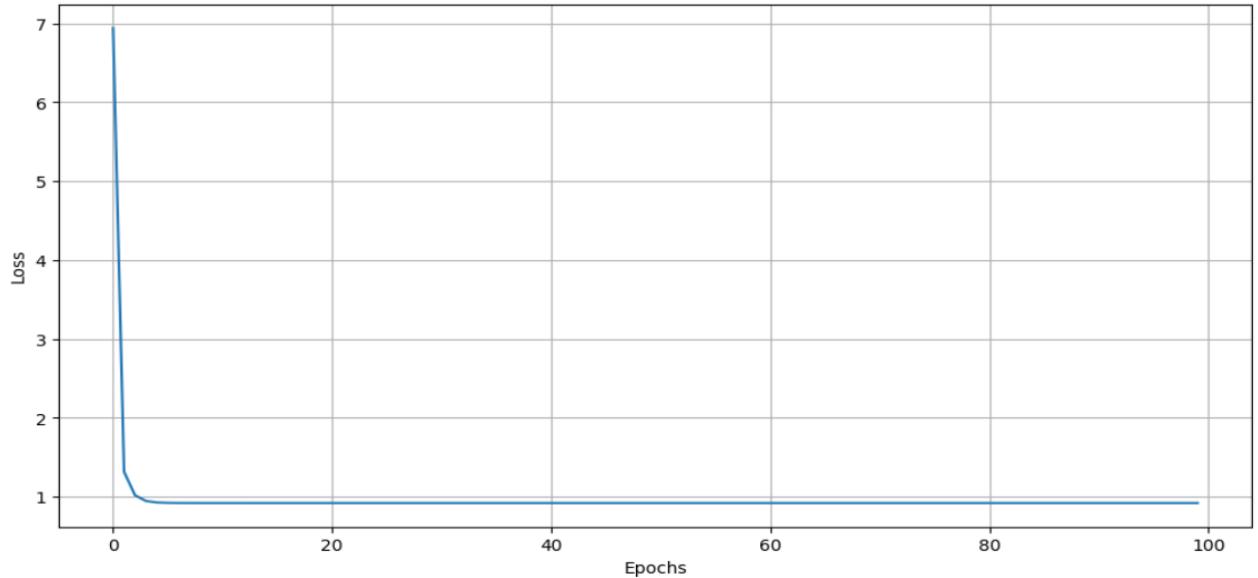


Stochastic Gradient Descent (SGD)

```
# Plot the loss curve
plt.figure(figsize=(12, 6))
plt.plot(range(epochs), loss_history, label='Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Loss Curve for Stochastic Gradient Descent')
plt.grid()
plt.show()
```

Learned Weights (W): [1.28056599 2.51443097 3.64171733]
Learned Bias (b): [2.81177659]

Loss Curve for Stochastic Gradient Descent

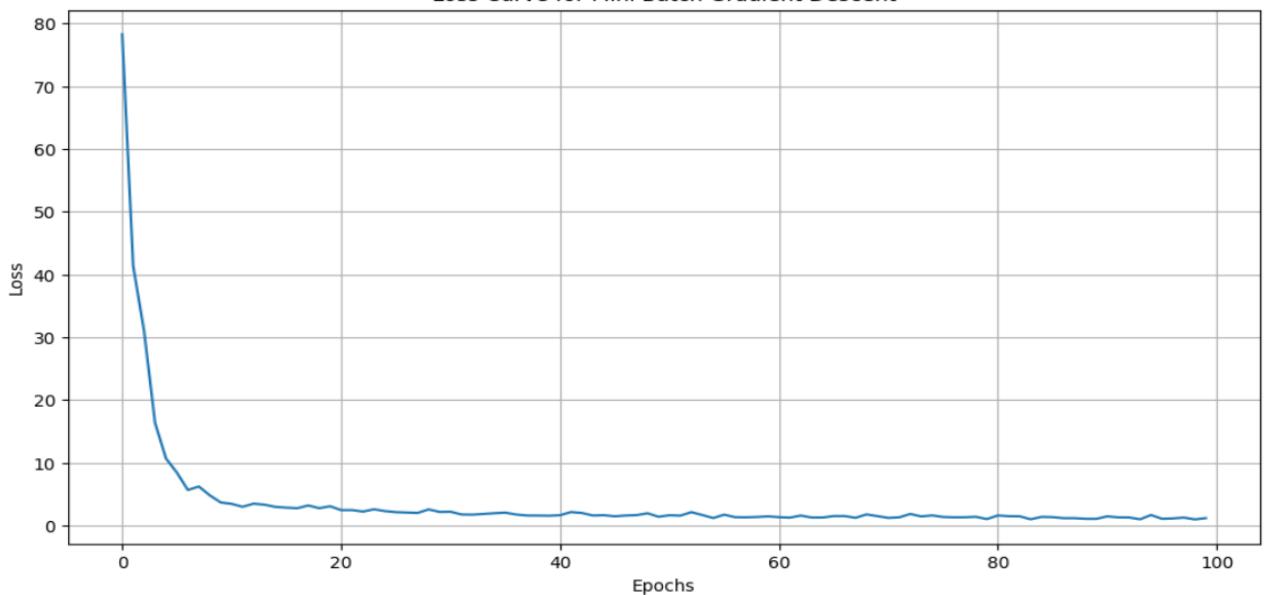


Mini-Batch Gradient Descent

```
# Plot the loss curve
plt.figure(figsize=(12, 6))
plt.plot(range(epochs), loss_history, label='Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Loss Curve for Mini-Batch Gradient Descent')
plt.grid()
plt.show()
```

Learned Weights (W): [1.466288 2.66537392 3.34412419]
Learned Bias (b): [2.89975528]

Loss Curve for Mini-Batch Gradient Descent



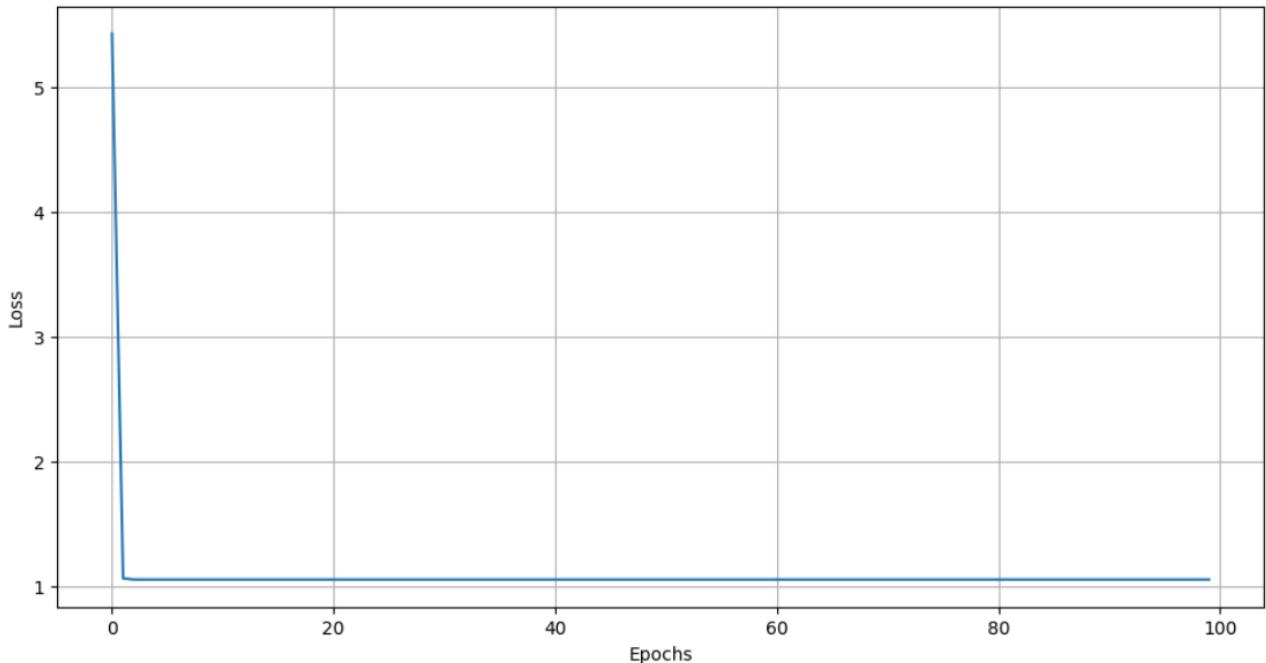


Momentum Gradient Descent

```
# Plot the loss curve
plt.figure(figsize=(12, 6))
plt.plot(range(epochs), loss_history, label='Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Loss Curve for Momentum Gradient Descent')
plt.grid()
plt.show()
```

Learned Weights (W): [1.53135691 2.41461102 3.45436068]
Learned Bias (b): [2.82780426]

Loss Curve for Momentum Gradient Descent





**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
(ARTIFICIAL INTELLIGENCE & MACHINE LEARNING)**

B.E / SEM VII / REV 2019 'C SCHEME' / CSE-(AI&ML)
Academic Year: 2023-24

NAME	PRATHAMESH S. CHIKANKAR
BRANCH	CSE-(AI&ML)
ROLL NO.	
SUBJECT	DEEP LEARNING LAB
COURSE CODE	CSL701
PRACTICAL NO.	
DOP	
DOS	



Output :-

DLL_EXP4.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
[9]  import numpy as np

# Define a simple dataset for XOR
X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])
y = np.array([[0], [1], [1], [0]])

# Define the neural network architecture
input_size, hidden_size1, hidden_size2, output_size = 2, 4, 3, 1
learning_rate, epochs = 0.1, 10000

# Initialize weights and biases
weights = [np.random.uniform(size=(input_size, hidden_size1)),
           np.random.uniform(size=(hidden_size1, hidden_size2)),
           np.random.uniform(size=(hidden_size2, output_size))]
biases = [np.zeros((1, hidden_size1)),
          np.zeros((1, hidden_size2)),
          np.zeros((1, output_size))]

# Sigmoid activation function and its derivative
sigmoid = lambda x: 1 / (1 + np.exp(-x))
sigmoid_derivative = lambda x: x * (1 - x)

# Training loop
for epoch in range(epochs):
    # Forward pass
    layers = [X]
    for w, b in zip(weights, biases):
        layers.append(sigmoid(np.dot(layers[-1], w) + b))
    output = layers[-1]

    # Backpropagation
    delta = [y - output]
    for i in range(len(weights) - 1, 0, -1):
        delta.append(delta[-1].dot(weights[i].T) * sigmoid_derivative(layers[i]))

    # Update weights and biases
    for i in range(len(weights)):
        weights[i] += layers[i].dot(delta[-1 - i]) * learning_rate
        biases[i] += np.sum(delta[-1 - i], axis=0, keepdims=True) * learning_rate

    if epoch % 1000 == 0:
        loss = np.mean(0.5 * (y - output) ** 2)
        print(f"Epoch {epoch}: Loss {loss}")

# Print final predictions
print("Final Predictions:")
print(output)
```

Epoch 0: Loss 0.14173703603991358
Epoch 1000: Loss 0.123212508461408
Epoch 2000: Loss 0.08778770918127633
Epoch 3000: Loss 0.07191910297219416
Epoch 4000: Loss 0.00022262241418865048
Epoch 5000: Loss 4.912702322947456e-05
Epoch 6000: Loss 2.0789432111937717e-05
Epoch 7000: Loss 1.1392281376228497e-05
Epoch 8000: Loss 7.174674996616704e-06
Epoch 9000: Loss 4.928971892945479e-06
Final Predictions:
[[0.00407558]
[0.99782183]
[0.99781978]
[0.00162783]]



Output :-

```
[1] import numpy as np

# Define the neural network architecture
input_size = 2
hidden_size1 = 3
hidden_size2 = 2
output_size = 1

# Initialize weights and biases for each layer
weights_input_hidden1 = np.random.randn(input_size, hidden_size1)
biases_hidden1 = np.zeros((1, hidden_size1))

weights_hidden1_hidden2 = np.random.randn(hidden_size1, hidden_size2)
biases_hidden2 = np.zeros((1, hidden_size2))

weights_hidden2_output = np.random.randn(hidden_size2, output_size)
biases_output = np.zeros((1, output_size))

# Sample input data
X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])
# Sample target output
y = np.array([[0], [1], [1], [0]])

# Define the learning rate
learning_rate = 0.01

# Training loop
epochs = 10000
for epoch in range(epochs):
    # Forward pass
    hidden1_input = np.dot(X, weights_input_hidden1) + biases_hidden1
    hidden1_output = 1 / (1 + np.exp(-hidden1_input))

    hidden2_input = np.dot(hidden1_output, weights_hidden1_hidden2) + biases_hidden2
    hidden2_output = 1 / (1 + np.exp(-hidden2_input))

    output = np.dot(hidden2_output, weights_hidden2_output) + biases_output

    # Calculate the loss
    loss = 0.5 * np.mean((output - y) ** 2)

    # Backpropagation
    dloss_doutput = output - y
    doutput_dhidden2_output = hidden2_output * (1 - hidden2_output)
    dloss_dhidden2_output = dloss_doutput.dot(weights_hidden2_output.T) * doutput_dhidden2_output

    dhidden2_output_dhidden1_output = hidden1_output * (1 - hidden1_output)
    dloss_dhidden1_output = dloss_dhidden2_output.dot(weights_hidden1_hidden2.T) * dhidden2_output_dhidden1_output

    # Update weights and biases
    weights_hidden2_output -= learning_rate * hidden2_output.T.dot(dloss_doutput)
    biases_output -= learning_rate * np.sum(dloss_doutput, axis=0, keepdims=True)

    weights_hidden1_hidden2 -= learning_rate * hidden1_output.T.dot(dloss_dhidden2_output)
    biases_hidden2 -= learning_rate * np.sum(dloss_dhidden2_output, axis=0, keepdims=True)

    weights_input_hidden1 -= learning_rate * X.T.dot(dloss_dhidden1_output)
    biases_hidden1 -= learning_rate * np.sum(dloss_dhidden1_output, axis=0, keepdims=True)

    if epoch % 1000 == 0:
        print(f'Epoch {epoch}, Loss: {loss}')
```



```
# Print final weights and biases
print("Final weights and biases:")
print("Weights Input to Hidden Layer 1:")
print(weights_input_hidden1)
print("Biases Hidden Layer 1:")
print(biases_hidden1)
print("Weights Hidden Layer 1 to Hidden Layer 2:")
print(weights_hidden1_hidden2)
print("Biases Hidden Layer 2:")
print(biases_hidden2)
print("Weights Hidden Layer 2 to Output Layer:")
print(weights_hidden2_output)
print("Biases Output Layer:")
print(biases_output)

print("Final Predictions")
print(output)
```

```
Epoch 0, Loss: 0.23812480513179063
Epoch 1000, Loss: 0.12593376620949448
Epoch 2000, Loss: 0.12568081295462774
Epoch 3000, Loss: 0.125501159986058
Epoch 4000, Loss: 0.12536576775687908
Epoch 5000, Loss: 0.12525762591421308
Epoch 6000, Loss: 0.12516583328803305
Epoch 7000, Loss: 0.12508264353810714
Epoch 8000, Loss: 0.1250018043609171
Epoch 9000, Loss: 0.12491744872159947
Final weights and biases:
Weights Input to Hidden Layer 1:
[[-0.67151782  0.75026953  0.39560969]
 [ 0.41462775 -0.21458599  0.91839141]]
Biases Hidden Layer 1:
[[-0.02962185 -0.04466992 -0.12989401]]
Weights Hidden Layer 1 to Hidden Layer 2:
[[-1.22731923 -1.82450893]
 [ 0.74629927  0.95909885]
 [-0.26582114  0.13230438]]
Biases Hidden Layer 2:
[[-0.21635442  0.19683121]]
Weights Hidden Layer 2 to Output Layer:
[[ 0.78343265]
 [-0.49679063]]
Biases Output Layer:
[[0.4538795]]
Final Predictions
[[0.50782372]
 [0.49428846]
 [0.50578433]
 [0.49061973]]
```



**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
(ARTIFICIAL INTELLIGENCE & MACHINE LEARNING)**

B.E / SEM VII / REV 2019 'C SCHEME' / CSE-(AI&ML)
Academic Year: 2023-24

NAME	PRATHAMESH S. CHIKANKAR
BRANCH	CSE-(AI&ML)
ROLL NO.	
SUBJECT	DEEP LEARNING LAB
COURSE CODE	CSL701
PRACTICAL NO.	
DOP	
DOS	



Output:

```
✓ [1] import keras
    from keras import layers
{x}
    # This is the size of our encoded representations
    encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats
    # This is our input image
    input_img = keras.Input(shape=(784,))

    # "encoded" is the encoded representation of the input
    encoded = layers.Dense(encoding_dim, activation='relu')(input_img)

    # "decoded" is the lossy reconstruction of the input
    decoded = layers.Dense(784, activation='sigmoid')(encoded)

    # This model maps an input to its reconstruction
    autoencoder = keras.Model(input_img, decoded)

    # This model maps an input to its encoded representation
    encoder = keras.Model(input_img, encoded)

    # This is our encoded (32-dimensional) input
    encoded_input = keras.Input(shape=(encoding_dim,))

    # Retrieve the last layer of the autoencoder model
    decoder_layer = autoencoder.layers[-1]

    # Create the decoder model
    decoder = keras.Model(encoded_input, decoder_layer(encoded_input))
    autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

✓ [2] from keras.datasets import mnist
    import numpy as np
    (x_train, _), (x_test, _) = mnist.load_data()

    Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
    11490434/11490434 [=====] - 1s 0us/step

✓ [3] x_train = x_train.astype('float32') / 255.
    x_test = x_test.astype('float32') / 255.
    x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
    x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))
    print(x_train.shape)
    print(x_test.shape)

    (60000, 784)
    (10000, 784)

✓ [4] autoencoder.fit(x_train, x_train,
                      epochs=50,
                      batch_size=256,
                      shuffle=True,
                      validation_data=(x_test, x_test))
```



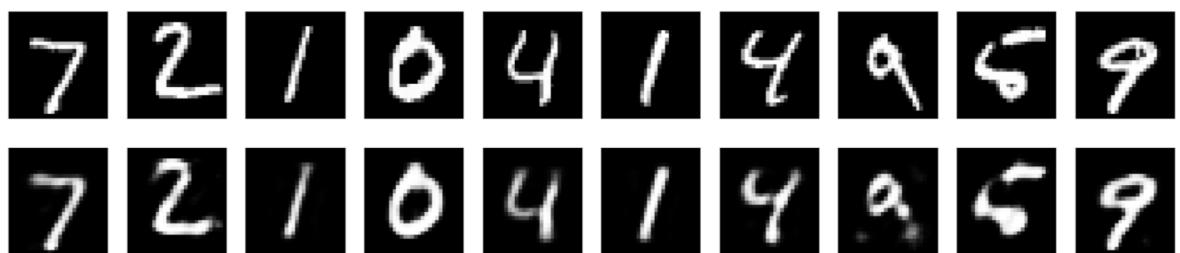
```
Epoch 1/50
235/235 [=====] - 6s 21ms/step - loss: 0.2768 - val_loss: 0.1897
Epoch 2/50
235/235 [=====] - 3s 11ms/step - loss: 0.1708 - val_loss: 0.1542
Epoch 3/50
235/235 [=====] - 3s 11ms/step - loss: 0.1453 - val_loss: 0.1353
Epoch 4/50
235/235 [=====] - 3s 11ms/step - loss: 0.1292 - val_loss: 0.1214
Epoch 5/50
235/235 [=====] - 3s 11ms/step - loss: 0.1182 - val_loss: 0.1120

Epoch 4/50
235/235 [=====] - 3s 13ms/step - loss: 0.0927 - val_loss: 0.0916
Epoch 48/50
235/235 [=====] - 3s 11ms/step - loss: 0.0927 - val_loss: 0.0915
Epoch 49/50
235/235 [=====] - 3s 11ms/step - loss: 0.0927 - val_loss: 0.0915
Epoch 50/50
235/235 [=====] - 3s 11ms/step - loss: 0.0927 - val_loss: 0.0915
<keras.callbacks.History at 0x7c36bd54ca60>
```

```
[5] # Encode and decode some digits
# Note that we take them from the *test* set
encoded_imgs = encoder.predict(x_test)
decoded_imgs = decoder.predict(encoded_imgs)
```

```
313/313 [=====] - 0s 1ms/step
313/313 [=====] - 0s 1ms/step
```

```
[6] # Use Matplotlib (don't ask)
import matplotlib.pyplot as plt
n = 10 # How many digits we will display
plt.figure(figsize=(20, 4))
for i in range(n):
    # Display original
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
    # Display reconstruction
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()
```





**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
(ARTIFICIAL INTELLIGENCE & MACHINE LEARNING)**

B.E / SEM VII / REV 2019 'C SCHEME' / CSE-(AI&ML)
Academic Year: 2023-24

NAME	PRATHAMESH S. CHIKANKAR
BRANCH	CSE-(AI&ML)
ROLL NO.	
SUBJECT	DEEP LEARNING LAB
COURSE CODE	CSL701
PRACTICAL NO.	
DOP	
DOS	

Output:

DL_EXP5-denoising.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

Load the necessary Libraries

```
[{x} 3s] [1] import numpy
      import matplotlib.pyplot as plt
      from keras.models import Sequential
      from keras.layers import Dense
      from keras.datasets import mnist
```

Load Dataset in Numpy Format

```
[{x} 0s] [2] (X_train, y_train), (X_test, y_test) = mnist.load_data()

      Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
      11490434/11490434 [=====] - 0s 0us/step
```

```
[{x} 0s] [3] X_train.shape
      (60000, 28, 28)
```

```
[{x} 0s] [4] X_test.shape
      (10000, 28, 28)
```

Plot Images as a Grey Scale Image

```
[{x} 1s] [5] plt.subplot(221)
      plt.imshow(X_train[0], cmap=plt.get_cmap('gray'))
      plt.subplot(222)
      plt.imshow(X_train[1], cmap=plt.get_cmap('gray'))
      plt.subplot(223)
      plt.imshow(X_train[2], cmap=plt.get_cmap('gray'))
      plt.subplot(224)
      plt.imshow(X_train[3], cmap=plt.get_cmap('gray'))
      # show the plot
      plt.show()
```



Formatting Data for Keras

```
✓ 0s [6] num_pixels = X_train.shape[1] * X_train.shape[2]
      X_train = X_train.reshape(X_train.shape[0], num_pixels).astype('float32')
      X_test = X_test.reshape(X_test.shape[0], num_pixels).astype('float32')
      X_train = X_train / 255
      X_test = X_test / 255

✓ 0s [7] X_train.shape
      (60000, 784)

✓ 0s [8] X_test.shape
      (10000, 784)
```

Adding Noise to Images

```
✓ 1s [9] noise_factor = 0.2
      x_train_noisy = X_train + noise_factor * numpy.random.normal(loc=0.0, scale=1.0, size=X_train.shape)
      x_test_noisy = X_test + noise_factor * numpy.random.normal(loc=0.0, scale=1.0, size=X_test.shape)
      x_train_noisy = numpy.clip(x_train_noisy, 0., 1.)
      x_test_noisy = numpy.clip(x_test_noisy, 0., 1.)
```

Defining an Encoder-Decoder network

```
✓ 1s [10] # create model
      model = Sequential()
      model.add(Dense(500, input_dim=num_pixels, activation='relu'))
      model.add(Dense(300, activation='relu'))
      model.add(Dense(100, activation='relu'))
      model.add(Dense(300, activation='relu'))
      model.add(Dense(500, activation='relu'))
      model.add(Dense(784, activation='sigmoid'))
```

Compiling the model

```
✓ 0s [11] # Compile the model
      model.compile(loss='mean_squared_error', optimizer='adam')
```

Training or Fitting the model

```
✓ 21s [12] # Training model
      model.fit(x_train_noisy, X_train, validation_data=(x_test_noisy, X_test), epochs=2, batch_size=200)

      Epoch 1/2
      300/300 [=====] - 10s 29ms/step - loss: 0.0422 - val_loss: 0.0203
      Epoch 2/2
      300/300 [=====] - 11s 35ms/step - loss: 0.0169 - val_loss: 0.0139
      <keras.src.callbacks.History at 0x793b0027d270>
```

Evaluating the model

```

✓ [13] # Final evaluation of the model
pred = model.predict(x_test_noisy)
pred.shape

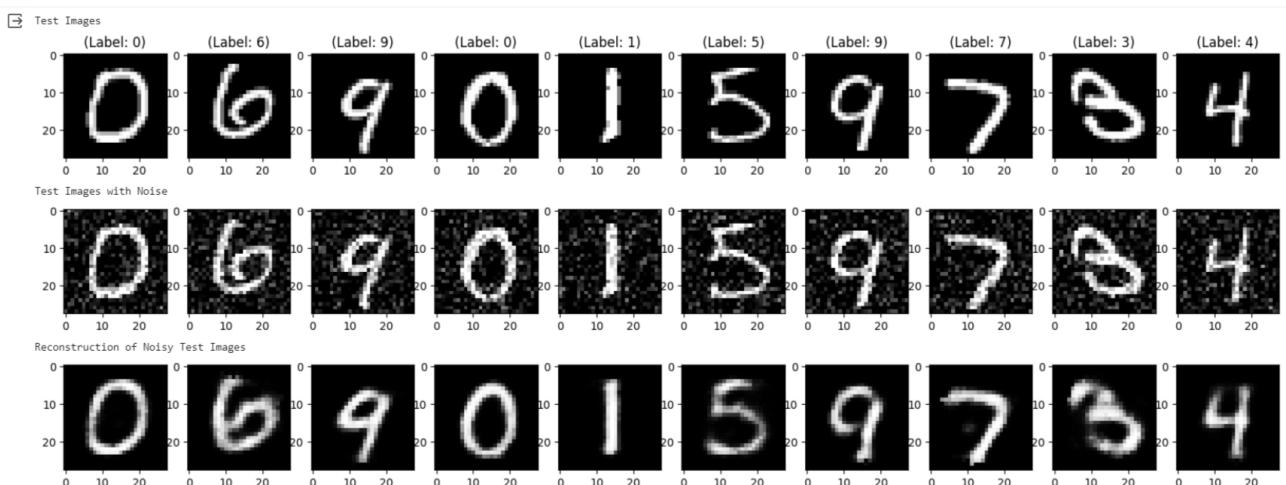
313/313 [=====] - 1s 4ms/step
(10000, 784)

✓ [14] X_test.shape

(10000, 784)

✓ [15] X_test = numpy.reshape(X_test, (10000,28,28)) *255
pred = numpy.reshape(pred, (10000,28,28)) *255
x_test_noisy = numpy.reshape(x_test_noisy, (-1,28,28)) *255
plt.figure(figsize=(20, 4))
print("Test Images")
for i in range(10,20,1):
    plt.subplot(2, 10, i+1)
    plt.imshow(X_test[i,:,:], cmap='gray')
    curr_lbl = y_test[i]
    plt.title("(Label: " + str(curr_lbl) + ")")
plt.show()
plt.figure(figsize=(20, 4))
print("Test Images with Noise")
for i in range(10,20,1):
    plt.subplot(2, 10, i+1)
    plt.imshow(x_test_noisy[i,:,:], cmap='gray')
plt.show()
plt.figure(figsize=(20, 4))
print("Reconstruction of Noisy Test Images")
for i in range(10,20,1):
    plt.subplot(2, 10, i+1)
    plt.imshow(pred[i,:,:], cmap='gray')
plt.show()

```





**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
(ARTIFICIAL INTELLIGENCE & MACHINE LEARNING)**

B.E / SEM VII / REV 2019 'C SCHEME' / CSE-(AI&ML)
Academic Year: 2023-24

NAME	PRATHAMESH S. CHIKANKAR
BRANCH	CSE-(AI&ML)
ROLL NO.	
SUBJECT	DEEP LEARNING LAB
COURSE CODE	CSL701
PRACTICAL NO.	
DOP	
DOS	

Output:

▼ Import TensorFlow

```

✓ 1s [2] import tensorflow as tf
    from tensorflow.keras import datasets, layers, models
    import matplotlib.pyplot as plt

✓ 18s [3] (train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()
    # Normalize pixel values to be between 0 and 1
    train_images, test_images = train_images / 255.0, test_images / 255.0

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 [=====] - 14s 0us/step

```

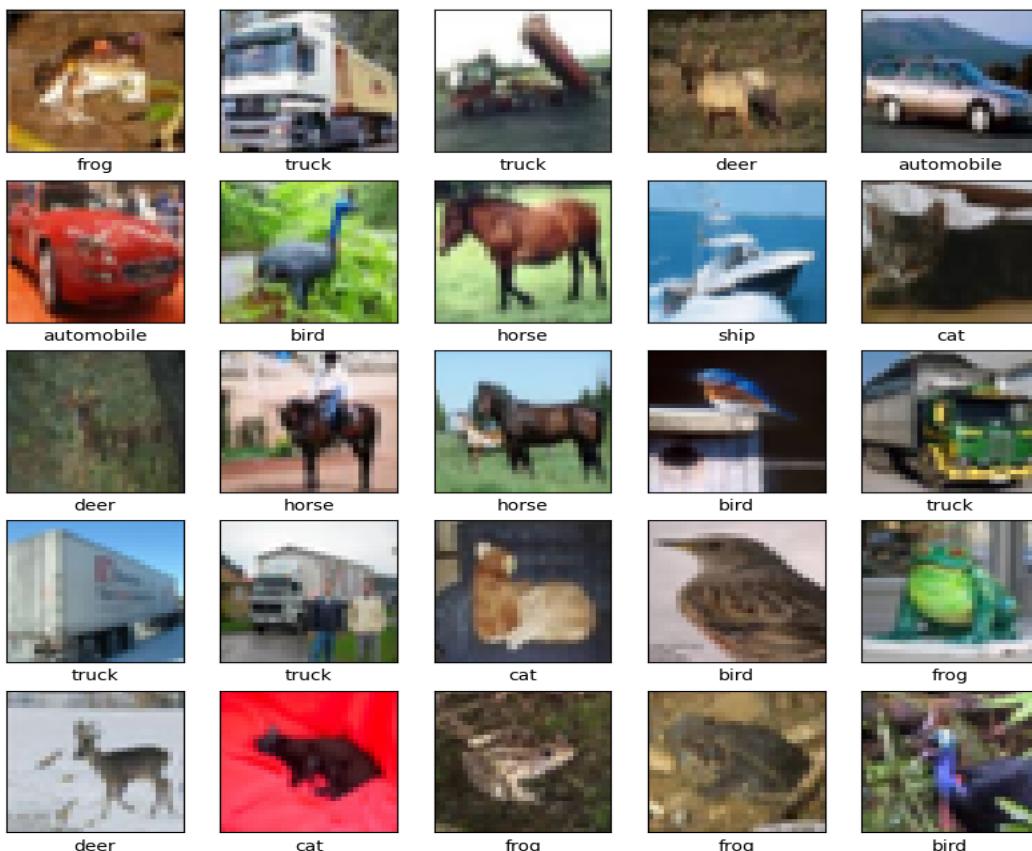
▼ Verify the data

```

✓ 2s [4] class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
                      'dog', 'frog', 'horse', 'ship', 'truck']

    plt.figure(figsize=(10,10))
    for i in range(25):
        plt.subplot(5,5,i+1)
        plt.xticks([])
        plt.yticks([])
        plt.grid(False)
        plt.imshow(train_images[i])
        # The CIFAR labels happen to be arrays,
        # which is why you need the extra index
        plt.xlabel(class_names[train_labels[i][0]])
    plt.show()

```





▼ Create the convolutional base

```
[5] model = models.Sequential()  
    model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))  
    model.add(layers.MaxPooling2D((2, 2)))  
    model.add(layers.Conv2D(64, (3, 3), activation='relu'))  
    model.add(layers.MaxPooling2D((2, 2)))  
    model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

Let's display the architecture of your model so far:

```
[6] model.summary()  
  
Model: "sequential"  
-----  
Layer (type)          Output Shape         Param #  
=====-----  
conv2d (Conv2D)       (None, 30, 30, 32)     896  
max_pooling2d (MaxPooling2D) (None, 15, 15, 32) 0  
conv2d_1 (Conv2D)      (None, 13, 13, 64)    18496  
max_pooling2d_1 (MaxPooling2D) (None, 6, 6, 64) 0  
conv2d_2 (Conv2D)      (None, 4, 4, 64)     36928  
=====  
Total params: 56320 (220.00 KB)  
Trainable params: 56320 (220.00 KB)  
Non-trainable params: 0 (0.00 Byte)
```

```
[7] model.add(layers.Flatten())  
model.add(layers.Dense(64, activation='relu'))  
model.add(layers.Dense(10))
```

Here's the complete architecture of model:

```
[8] model.summary()  
  
Model: "sequential"  
-----  
Layer (type)          Output Shape         Param #  
=====-----  
conv2d (Conv2D)       (None, 30, 30, 32)     896  
max_pooling2d (MaxPooling2D) (None, 15, 15, 32) 0  
conv2d_1 (Conv2D)      (None, 13, 13, 64)    18496  
max_pooling2d_1 (MaxPooling2D) (None, 6, 6, 64) 0  
conv2d_2 (Conv2D)      (None, 4, 4, 64)     36928  
flatten (Flatten)     (None, 1024)           0  
dense (Dense)         (None, 64)            65600  
dense_1 (Dense)       (None, 10)             650  
=====  
Total params: 122570 (478.79 KB)  
Trainable params: 122570 (478.79 KB)  
Non-trainable params: 0 (0.00 Byte)
```



▼ Compile and train the model

```
✓ [9] model.compile(optimizer='adam',
                     loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                     metrics=['accuracy'])

history = model.fit(train_images, train_labels, epochs=10,
                     validation_data=(test_images, test_labels))

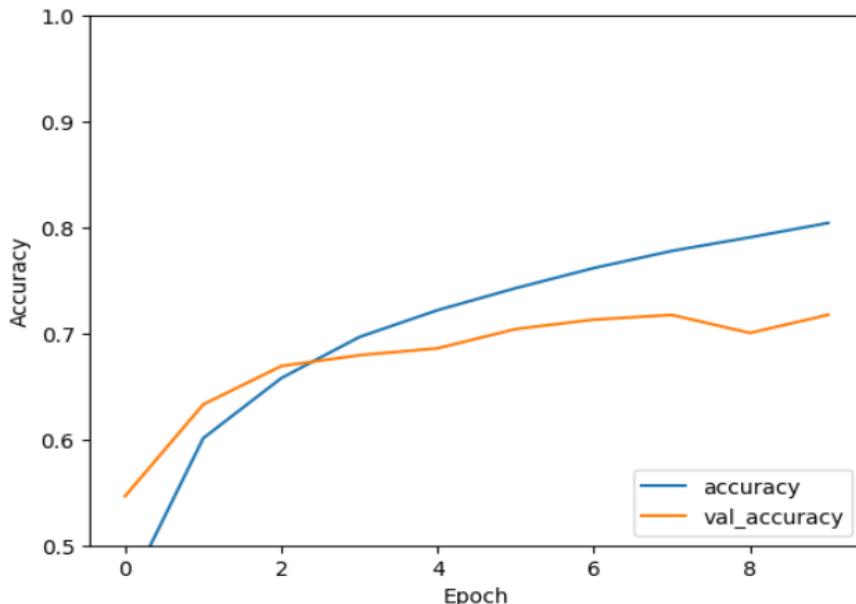
Epoch 1/10
1563/1563 [=====] - 23s 6ms/step - loss: 1.5035 - accuracy: 0.4517 - val_loss: 1.2647 - val_accuracy: 0.5465
Epoch 2/10
1563/1563 [=====] - 9s 6ms/step - loss: 1.1228 - accuracy: 0.6011 - val_loss: 1.0465 - val_accuracy: 0.6330
Epoch 3/10
1563/1563 [=====] - 8s 5ms/step - loss: 0.9661 - accuracy: 0.6579 - val_loss: 0.9493 - val_accuracy: 0.6693
Epoch 4/10
1563/1563 [=====] - 8s 5ms/step - loss: 0.8638 - accuracy: 0.6967 - val_loss: 0.9379 - val_accuracy: 0.6795
Epoch 5/10
1563/1563 [=====] - 9s 6ms/step - loss: 0.7929 - accuracy: 0.7220 - val_loss: 0.9127 - val_accuracy: 0.6860
Epoch 6/10
1563/1563 [=====] - 9s 6ms/step - loss: 0.7316 - accuracy: 0.7428 - val_loss: 0.8524 - val_accuracy: 0.7042
Epoch 7/10
1563/1563 [=====] - 8s 5ms/step - loss: 0.6782 - accuracy: 0.7618 - val_loss: 0.8613 - val_accuracy: 0.7131
Epoch 8/10
1563/1563 [=====] - 8s 5ms/step - loss: 0.6359 - accuracy: 0.7779 - val_loss: 0.8538 - val_accuracy: 0.7175
Epoch 9/10
1563/1563 [=====] - 9s 6ms/step - loss: 0.5925 - accuracy: 0.7908 - val_loss: 0.9277 - val_accuracy: 0.7005
Epoch 10/10
1563/1563 [=====] - 8s 5ms/step - loss: 0.5552 - accuracy: 0.8043 - val_loss: 0.8984 - val_accuracy: 0.7176
```

▼ Evaluate the model

```
✓ [10] plt.plot(history.history['accuracy'], label='accuracy')
        plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
        plt.xlabel('Epoch')
        plt.ylabel('Accuracy')
        plt.ylim([0.5, 1])
        plt.legend(loc='lower right')

test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)

313/313 - 1s - loss: 0.8984 - accuracy: 0.7176 - 687ms/epoch - 2ms/step
```



```
✓ [11] print(test_acc)

0.7175999879837036
```



**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
(ARTIFICIAL INTELLIGENCE & MACHINE LEARNING)**

B.E / SEM VII / REV 2019 'C SCHEME' / CSE-(AI&ML)
Academic Year: 2023-24

NAME	PRATHAMESH S. CHIKANKAR
BRANCH	CSE-(AI&ML)
ROLL NO.	
SUBJECT	DEEP LEARNING LAB
COURSE CODE	CSL701
PRACTICAL NO.	
DOP	
DOS	



Output:

```
!pip install scalecast --upgrade
```

```
!pip install tensorflow
```

🔍 ▾ Data Preprocessing

```
{x}
```

```
✓ [11] import pandas as pd
     import numpy as np
     import pickle
     import seaborn as sns
     import matplotlib.pyplot as plt
     from scalecast.Forecaster import Forecaster

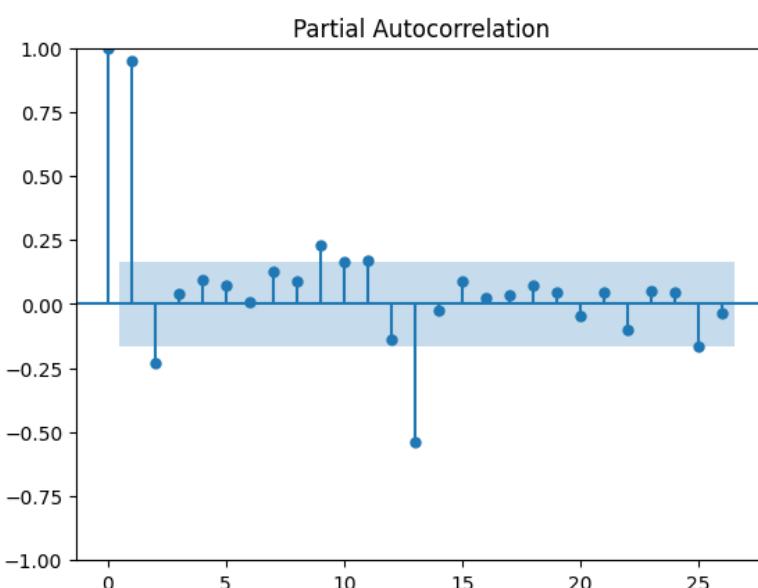
     df = pd.read_csv('AirPassengers.csv',parse_dates=['Month'])
```

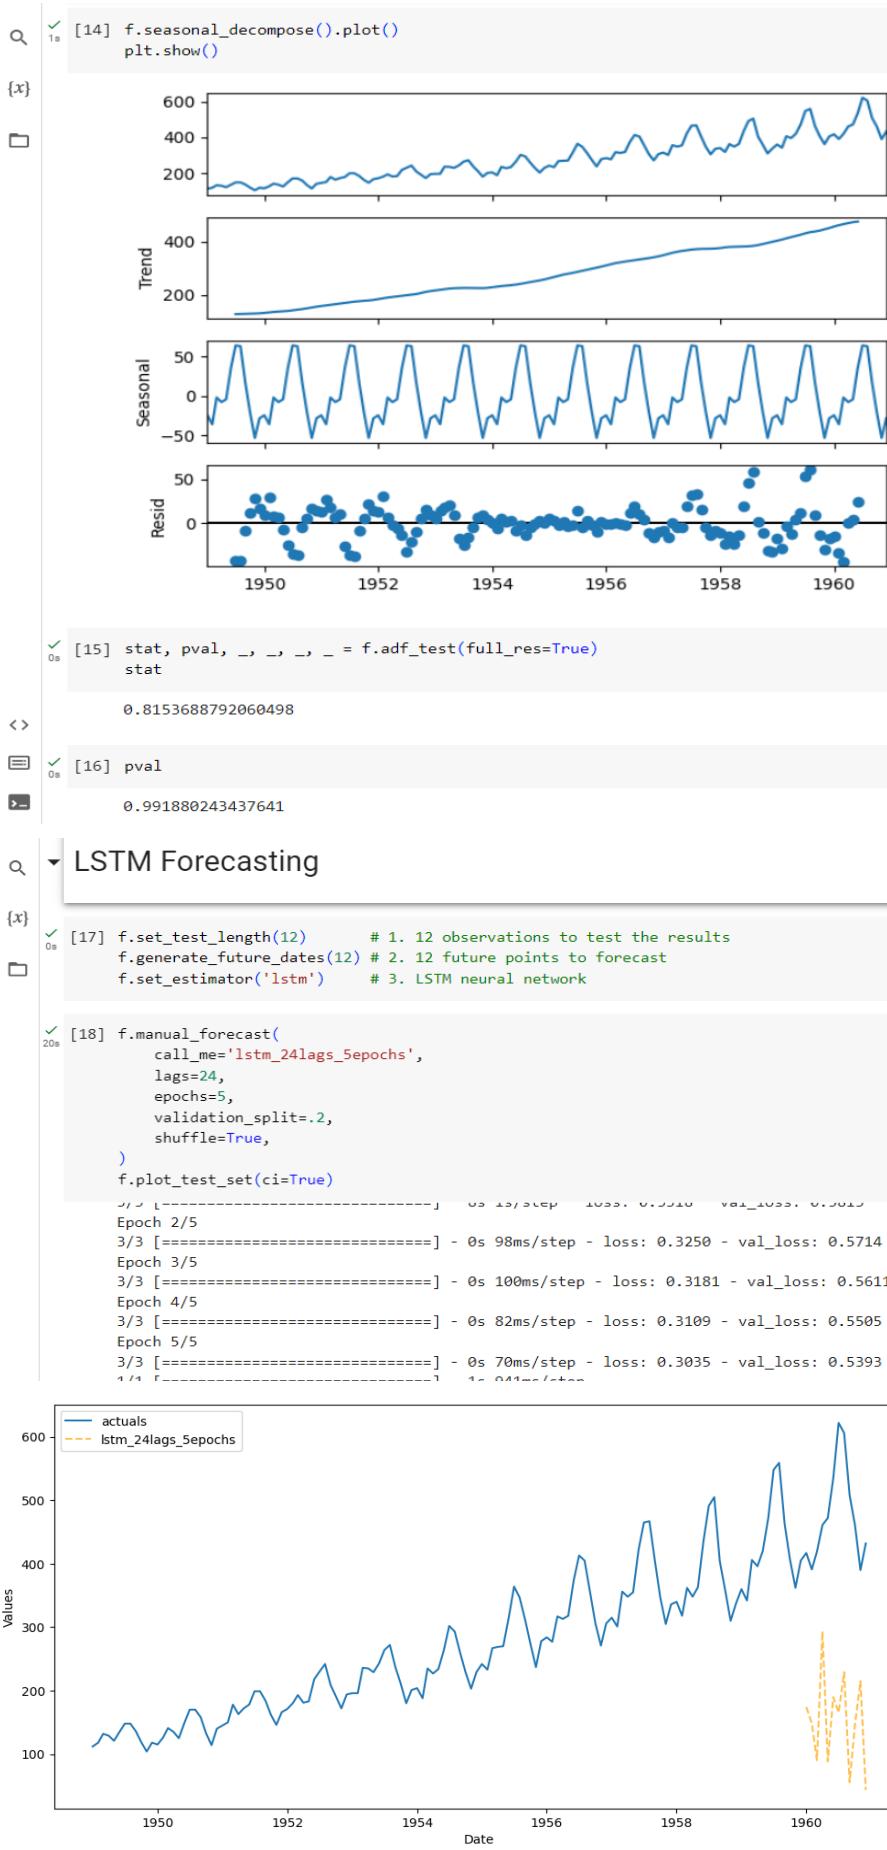
▾ Exploratory Data Analysis

```
✓ [12] f = Forecaster(y=df['#Passengers'],current_dates=df['Month'])
     f
```

```
Forecaster(
    DateStartActuals=1949-01-01T00:00:00.000000000
    DateEndActuals=1960-12-01T00:00:00.000000000
    Freq=MS
    N_actuals=144
    ForecastLength=0
    Xvars=[]
    TestLength=0
    ValidationMetric=rmse
    ForecastsEvaluated=[]
    CILevel=None
    CurrentEstimator=mlr
    GridsFile=Grids
)
```

```
✓ [13] f.plot_pacf(lags=26)
     plt.show()
```





```
[15] stat, pval, _, _, _, _ = f.adf_test(full_res=True)  
stat
```

0.8153688792060498

```
[16] pval
```

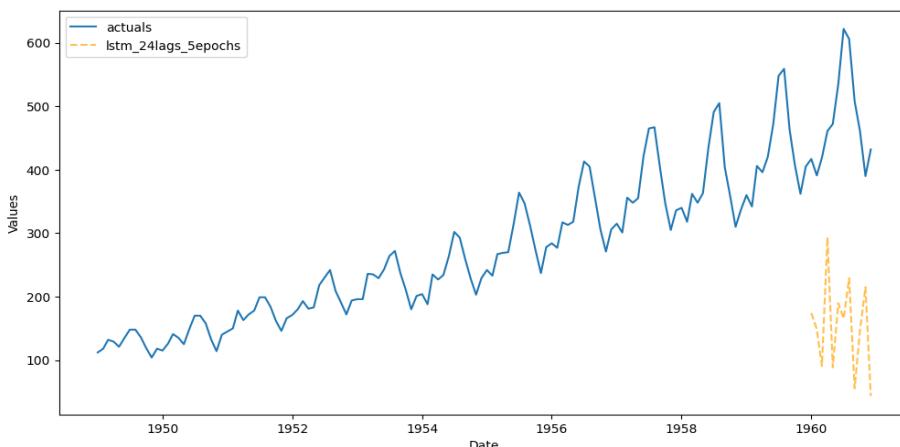
0.991880243437641

LSTM Forecasting

```
[17] f.set_test_length(12) # 1. 12 observations to test the results  
f.generate_future_dates(12) # 2. 12 future points to forecast  
f.set_estimator('lstm') # 3. LSTM neural network
```

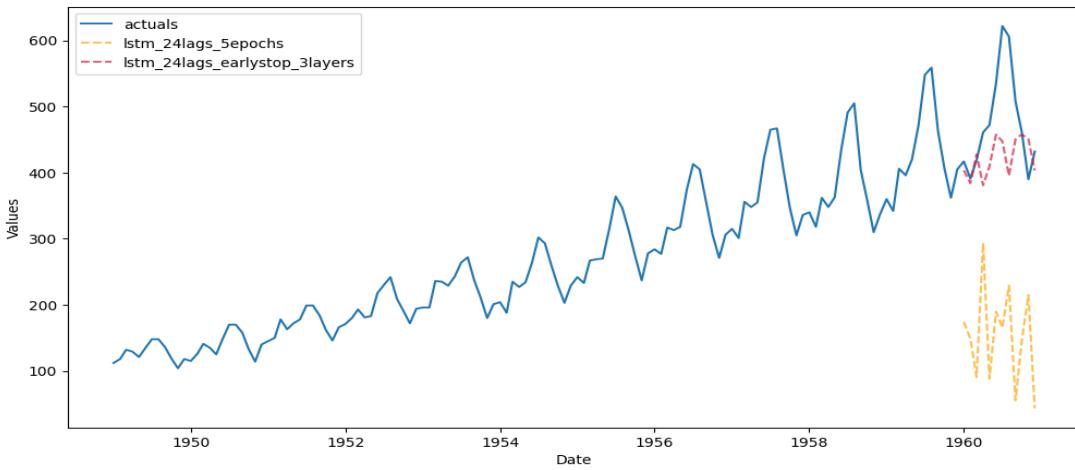
```
[18] f.manual_forecast(  
    call_me='lstm_24lags_5epochs',  
    lags=24,  
    epochs=5,  
    validation_split=.2,  
    shuffle=True,  
)  
f.plot_test_set(ci=True)
```

Epoch 1/5
Epoch 2/5
3/3 [=====] - 0s 98ms/step - loss: 0.3250 - val_loss: 0.5714
Epoch 3/5
3/3 [=====] - 0s 100ms/step - loss: 0.3181 - val_loss: 0.5611
Epoch 4/5
3/3 [=====] - 0s 82ms/step - loss: 0.3109 - val_loss: 0.5505
Epoch 5/5
3/3 [=====] - 0s 70ms/step - loss: 0.3035 - val_loss: 0.5393





```
{4} [19] from tensorflow.keras.callbacks import EarlyStopping  
f.manual_forecast(  
    call_me='lstm_24lags_earlystop_3layers',  
    lags=24,  
    epochs=25,  
    validation_split=.2,  
    shuffle=True,  
    callbacks=EarlyStopping(  
        monitor='val_loss',  
        patience=5,  
    ),  
    lstm_layer_sizes=(16,16,16),  
    dropout=(0,0,0),  
)  
  
f.plot_test_set(ci=True)  
  
Epoch 1/25  
3/3 [=====] - 10s 922ms/step - loss: 0.3553 - val_loss: 0.6296  
Epoch 2/25  
3/3 [=====] - 0s 63ms/step - loss: 0.3400 - val_loss: 0.6073  
Epoch 3/25  
3/3 [=====] - 0s 80ms/step - loss: 0.3245 - val_loss: 0.5831  
Epoch 4/25  
3/3 [=====] - 0s 86ms/step - loss: 0.3071 - val_loss: 0.5544  
Epoch 5/25  
3/3 [=====] - 0s 78ms/step - loss: 0.2859 - val_loss: 0.5174  
Epoch 6/25  
3/3 [=====] - 0s 68ms/step - loss: 0.2583 - val_loss: 0.4663  
Epoch 7/25  
3/3 [=====] - 0s 54ms/step - loss: 0.0824 - val_loss: 0.1556  
Epoch 24/25  
3/3 [=====] - 0s 50ms/step - loss: 0.0781 - val_loss: 0.1520  
Epoch 25/25  
3/3 [=====] - 0s 53ms/step - loss: 0.0749 - val_loss: 0.1470  
WARNING:tensorflow:5 out of the last 13 calls to <function Model.make_predict_function.<locals>.predict_function at 0x797362aa36d0>  
1/1 [=====] - 1s 1s/step  
4/4 [=====] - 0s 9ms/step  
/usr/local/lib/python3.10/dist-packages/scalecast/_utils.py:60: Warning: Confidence intervals not found for lstm_24lags_5epochs. To  
    warnings.warn()  
/usr/local/lib/python3.10/dist-packages/scalecast/_utils.py:60: Warning: Confidence intervals not found for lstm_24lags_earlystop_3.  
    warnings.warn()  
<Axes: xlabel='Date', ylabel='Values'>
```



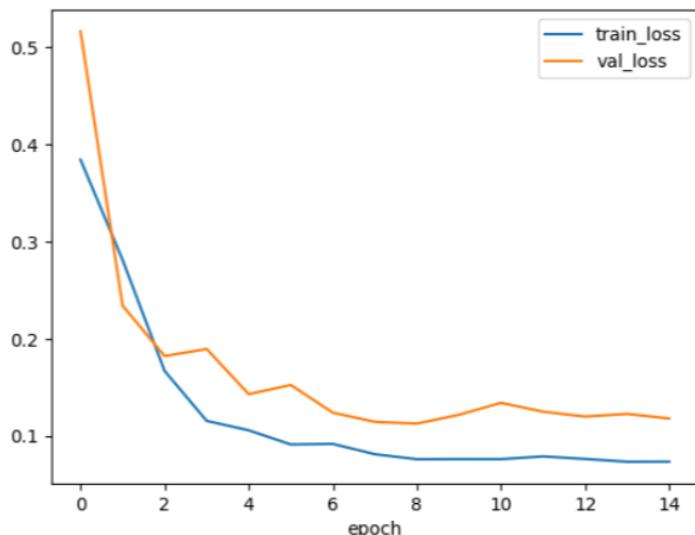
```
[21] f.manual_forecast(  
    call_me='lstm_best',  
    lags=36,  
    batch_size=32,  
    epochs=15,  
    validation_split=.2,  
    shuffle=True,  
    activation='tanh',  
    optimizer='Adam',  
    learning_rate=0.001,  
    lstm_layer_sizes=(72,)*4,  
    dropout=(0,)*4,  
    plot_loss=True  
)  
f.plot_test_set(order_by='TestSetMAPE', models='top_2', ci=True)
```

```

Epoch 1/15
3/3 [=====] - 11s 1s/step - loss: 0.3839 - val_loss: 0.5157
Epoch 2/15
3/3 [=====] - 0s 97ms/step - loss: 0.2809 - val_loss: 0.2339
Epoch 3/15
3/3 [=====] - 0s 97ms/step - loss: 0.1668 - val_loss: 0.1821
Epoch 4/15
3/3 [=====] - 0s 98ms/step - loss: 0.0735 - val_loss: 0.1226
Epoch 14/15
3/3 [=====] - 0s 98ms/step - loss: 0.0736 - val_loss: 0.1179
1/1 [=====] - 1s 1s/step

```

Istm model loss

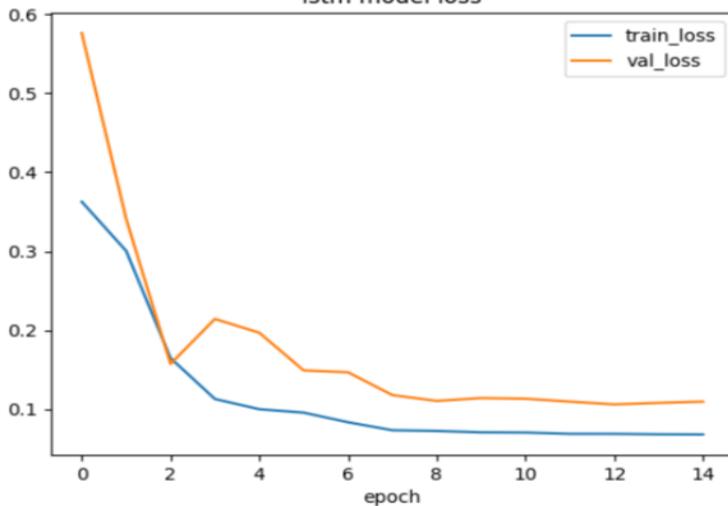


```

Epoch 1/15
3/3 [=====] - 10s 864ms/step - loss: 0.3627 - val_loss: 0.5764
Epoch 2/15
3/3 [=====] - 0s 106ms/step - loss: 0.3004 - val_loss: 0.3410
Epoch 3/15
3/3 [=====] - 0s 103ms/step - loss: 0.1646 - val_loss: 0.1572
Epoch 14/15
3/3 [=====] - 0s 112ms/step - loss: 0.0680 - val_loss: 0.1078
Epoch 15/15
3/3 [=====] - 0s 106ms/step - loss: 0.0678 - val_loss: 0.1094
1/1 [=====] - 2s 2s/step
3/3 [=====] - 0s 39ms/step

```

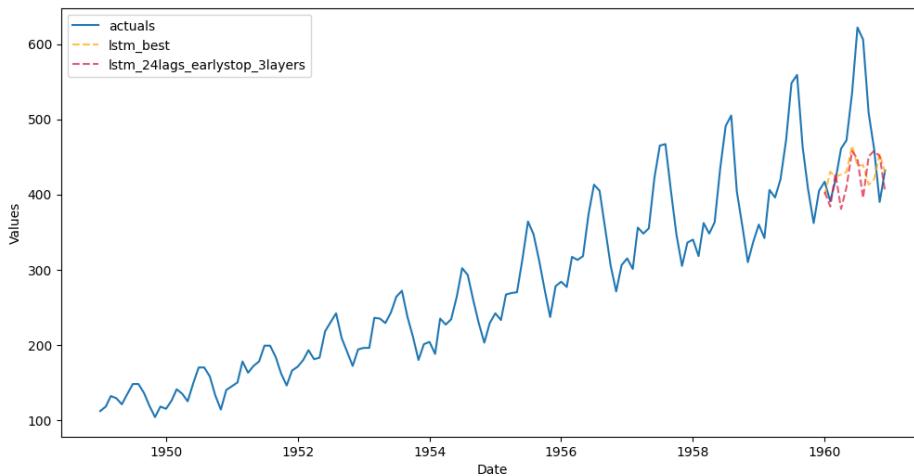
Istm model loss



```

/usr/local/lib/python3.10/dist-packages/scalecast/_utils.py:60: Warning: Confidence inte
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/scalecast/_utils.py:60: Warning: Confidence inte
    warnings.warn(
<Axes: xlabel='Date', ylabel='Values'>

```



{x} ▾ MLR Forecasting and Model Benchmarking

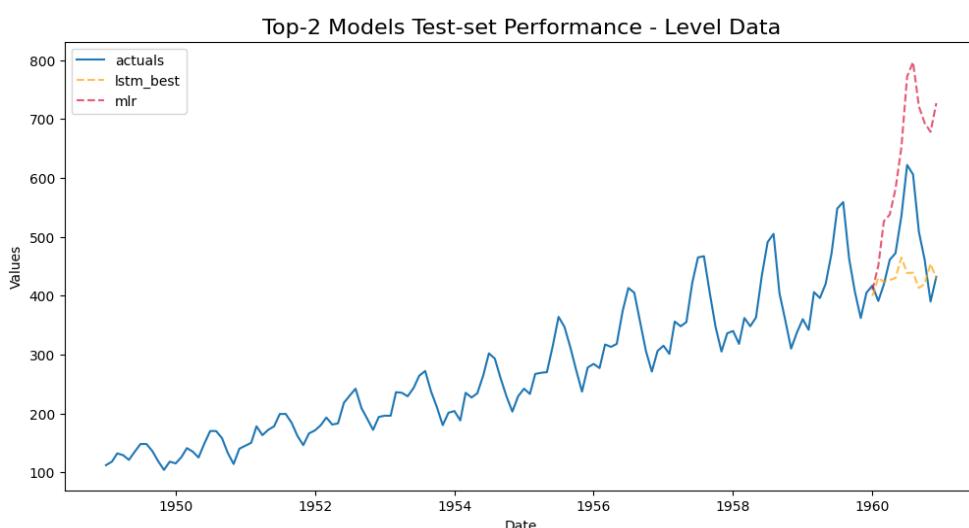
```
✓ [45] from scalecast.SeriesTransformer import SeriesTransformer
      transformer = SeriesTransformer(f)
      f = transformer.DiffTransform()

      f.add_ar_terms(24)
      f.add_seasonal_regressors('month','quarter',dummy=True)
      f.add_seasonal_regressors('year')
      f.add_time_trend()

✓ [46] f.set_estimator('mlr')
      f.manual_forecast()

      f = transformer.DiffRevert(
          exclude_models = [m for m in f.history if m != 'mlr']
      ) # exclude all lstm models from the revert

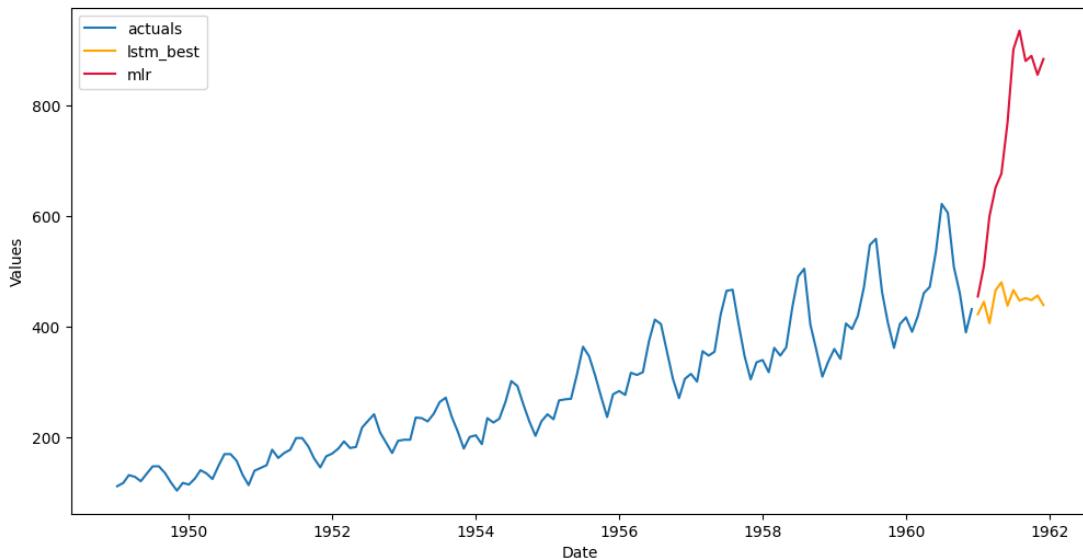
      f.plot_test_set(order_by='TestSetMAPE',models=['lstm_best','mlr'])
      plt.title('Top-2 Models Test-set Performance - Level Data',size=16)
      plt.show()
```





```
✓ [57] f.plot(models=['mlr', 'lstm_best'], order_by='TestSetMAPE')
```

```
<Axes: xlabel='Date', ylabel='Values'>
```



```
✓ [65] f.export('model_summaries',determine_best_by='InSampleRMSE')[['ModelNickname','InSampleR2','TestSetRMSE','TestSetMAPE','TestSetMAE','TestSetR2','best_model']]
```

	ModelNickname	InSampleR2	TestSetRMSE	TestSetMAPE	TestSetMAE	TestSetR2	best_model
0	mlr	0.970799	176.765019	0.328659	153.879150	-4.640586	True
1	lstm_best	0.704907	84.740822	0.122145	63.453901	-0.296334	False
2	lstm_24lags_earlystop_3layers	0.643474	90.941170	0.124324	65.348639	-0.492976	False
3	lstm_24lags_5epochs	-0.816249	336.169279	0.673811	322.878162	-19.400822	False



**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
(ARTIFICIAL INTELLIGENCE & MACHINE LEARNING)**

B.E / SEM VII / REV 2019 'C SCHEME' / CSE-(AI&ML)
Academic Year: 2023-24

NAME	SINGH SUDHAM DHARMENDRA
BRANCH	CSE-(AI&ML)
ROLL NO.	
SUBJECT	DEEP LEARNING LAB
COURSE CODE	CSL701
PRACTICAL NO.	
DOP	
DOS	



Output:

!pip install tensorflow

```
✓ [19] import numpy as np
       import tensorflow as tf
       from tensorflow import keras
{x}    from tensorflow.keras.models import Sequential
       from tensorflow.keras.layers import SimpleRNN, Dense, Embedding
       from tensorflow.keras.preprocessing.text import Tokenizer
       from tensorflow.keras.preprocessing.sequence import pad_sequences

# Sample data for character-level translation
input_texts = ['hello', 'world', 'good', 'bye', 'prathmesh', 'op', 'sudhamesh']
target_texts = ['olleh', 'dlrow', 'doog', 'eyb', 'hsemhtarp', 'po', 'hsemahdus']

# Tokenize input and target texts
tokenizer = Tokenizer(char_level=True)
tokenizer.fit_on_texts(input_texts + target_texts)
input_sequences = tokenizer.texts_to_sequences(input_texts)
target_sequences = tokenizer.texts_to_sequences(target_texts)

# Determine the maximum sequence length
max_seq_length = max(len(seq) for seq in input_sequences)

# Pad sequences to the maximum length
input_sequences = pad_sequences(input_sequences, maxlen=max_seq_length)
target_sequences = pad_sequences(target_sequences, maxlen=max_seq_length)

# Define the RNN model
model = Sequential()
model.add(Embedding(input_dim=len(tokenizer.word_index) + 1, output_dim=8, input_length=max_seq_length))
model.add(SimpleRNN(200, return_sequences=True))
model.add(Dense(len(tokenizer.word_index) + 1, activation='softmax'))

# Compile the model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy')

# Train the model
model.fit(input_sequences, target_sequences, epochs=50)

# Translate new input sequences
new_input_texts = ['prathmesh', 'sudhamesh']
new_input_sequences = tokenizer.texts_to_sequences(new_input_texts)
new_input_sequences = pad_sequences(new_input_sequences, maxlen=max_seq_length)
predicted_sequences = model.predict(new_input_sequences)
predicted_sequences = np.argmax(predicted_sequences, axis=-1)
for i, text in enumerate(new_input_texts):
    input_seq = new_input_sequences[i]
    predicted_seq = predicted_sequences[i]
    decoded_text = tokenizer.sequences_to_texts([predicted_seq])[0]
    print(f'Input: {text}, Predicted: {decoded_text}')

Epoch 1/50
1/1 [=====] - 2s 2s/step - loss: 2.8332
Epoch 2/50
1/1 [=====] - 0s 17ms/step - loss: 2.8066
Epoch 3/50
1/1 [=====] - 0s 13ms/step - loss: 2.7785
Epoch 4/50
Epoch 4/50
1/1 [=====] - 0s 17ms/step - loss: 1.0980
Epoch 48/50
1/1 [=====] - 0s 13ms/step - loss: 1.0764
Epoch 49/50
1/1 [=====] - 0s 16ms/step - loss: 1.0530
Epoch 50/50
1/1 [=====] - 0s 14ms/step - loss: 1.0269
1/1 [=====] - 0s 192ms/step
Input: prathmesh, Predicted: m h t a r p
Input: sudhamesh, Predicted: m a h d u s
```

Name :- PruthvrajeSH S. Chikankar.

ROLL NO. :- AJML11 BRANCH :- CSE - (AJML)

Year :- BE Subject :- Deep learning [DL]

Topic :- Assignment No. 01

Date :- August 23 SIGN :- PruthvrajeSH



1. what is Deep Learning?

- > Deep Learning is a subset of machine learning that focuses on artificial Neural Networks with multiple layers, also known as Deep Neural Networks. These networks are designed to automatically learn and extract hierarchical features from data.
- > Explain three different classes of deep networks.
- (*) Feedforward Neural Net (FFNN) :
 - 1) A FFNN is an artificial Neural Network where connections between the nodes are directed and do not form a cycle.
 - 2) As such it is different from its descendent: Recurrent Neural Net. The feed forward neural net was the first & simplest type of neural net device.
 - 3) In an artificial feed forward network (FFNN), information always move in one direction, i.e. it never goes backwards.
 - (*) Convolution Neural Net:
 - 1) A CNN (Convolution Neural Net) is a type of ANN used for major recognition. CNNs are a category of neural nets that have proven very effective in arrays such as image recognition and classification.
 - 2) CNNs are important tools for ML practitioners today. CNN is a deep learning neural net designed for processing structured arrays of data such as images.
 - (*) Recurrent Neural Net (RNN) :
 - 1) RNN neural net is a class of artificial Neural Net where connections between nodes form a directed graph along a temporal sequence.



This makes it to show temporal dynamic behaviour.

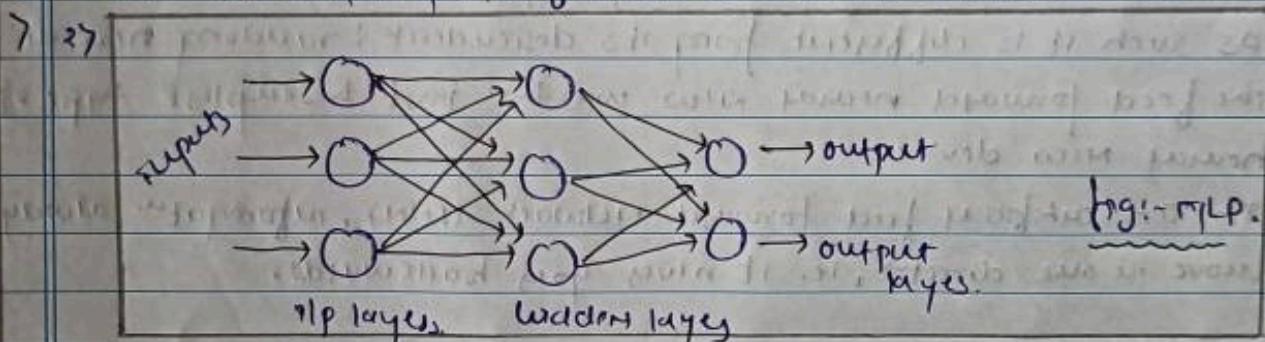
- It uses sequential data or time series data. These data are commonly used for ordinary or temporal problems, such as language translation, NLP, speech recognition & image captioning.

2. Explain multilayer perceptrons (MLPs) and represent power of MLPs.

- MLP is connected dense layers, and that transforms any input dimensions to desired dimensions.

A MLP is a neural net that has multiple layers.

MLP are so called, referred as "vanilla" Neural Network.



- In above diagram of MLP, there are three I/P & thus three O/P nodes and the hidden layer has three nodes.
- The O/P layer given two O/P, hence there are two O/P nodes. The nodes in the I/P layer take I/P & forward it for further process.
- The hidden layers process the I/P & passes it to the O/P layer.

- 3) Representational power! (Activation fn) -

- ^{upt} for the I/P nodes, each node is a neuron that uses a non-linear 'activation function'. MLP utility a chart based 'supervised learning' technique called back propagation for training.

- The two common activation function are Sigmoidal, and are given as:

$$y(x_i) = \tanh(x_i) \text{ and } y(x_i) = (1 + e^{-x_i}) = \frac{1}{1 + e^{-x_i}}$$



- The first is a 'hyperbolic tangent' that ranges from -1 to 1. While the other is 'logistic function' which is similar in shape but ranges from 0 to 1.
- Here y_i is the o/p & x_i is the weighted sum with node \uparrow of all p connect.
- alternate activation function have also been recognised, including the Rectified and softplus funn.

3.

Explain the working of multilayer feed forward NN & backpropagation in detail.

* Multilayer feed forward NN:

1) A multilayer feed forward NN (also known as a feed forward NN or FNN) works by passing data through multiple layers of neurons without any feedback loop.

2) Working:

- Input layer: the ip layer receives the raw data features.

- Hidden layers: There can be one or more hidden layers, each consisting of multiple neurons. Each neuron computes a weighted sum of its ip, applying an activation function (e.g.: sigmoid or ReLU) and passes the result to the next layer.

- Output layer: the O/p layer produces the final prediction or classification. The no. of neurons in this layer depends on the specific task (e.g., one neuron for binary classification, multiple neurons for multi-class classification).

- The weights and biases of neurons are initialized randomly, and the NN o/p is compared to the true target value using a loss function (e.g.: mean squared error for regression or cross-entropy for classification).



> * Back propagation:

- 1) Back propagation: It used to adjust the weights & bias to minimize the loss.
- 2) Working:
 - Forward pass: the input data is passed through the HW, and predictions are obtained.
 - Loss computation: the loss is calculated by comparing predict to the true labels.
 - backward pass (Backpropagation): Errors are propagated backward through the HW. The gradient of the loss with respect to the weights and bias are computed using the chain rule of calculus.
 - Weight update: the weight & bias are updated in the opposite direction of the gradients to minimize the loss. This is typically done using optimization algorithm like gradient descent.
 - The process is repeated iteratively until the model convergence to the set of weights & bias that produce accurate prediction.

4.

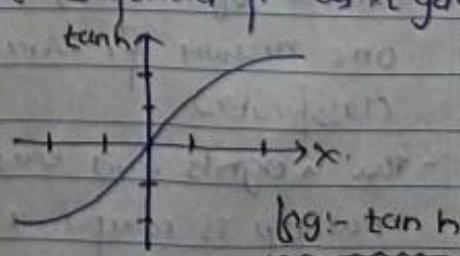
Explains different activation functions in detail,

1) \tanh fn: Used for neural Netwo-

$$f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} = \tanh(x)$$

by \tanh fn because preferred over the sigmoid fn as it gives better performance for multi-layer Neural Netwo.

- But it doesn't solve vanishing gradient problem.



2) Logistic activation fn: Also known as sigmoid fn.



> This fn is used in back-propagation nets and can be of two types.

- ① Binary sigmoid fn:

Also known as unipolar sigmoid fn or logistic sigmoid fn.

Defined as $f(x) = \frac{1}{1+e^{-\lambda x}}$

- Range from 0 to 1.

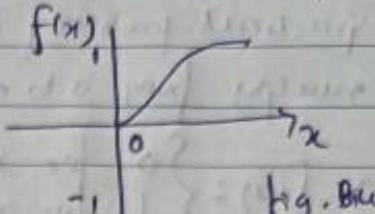
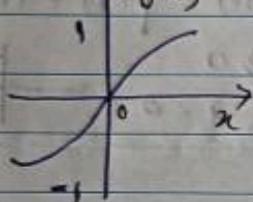


fig. Binary sigmoid fn.

$\therefore \lambda \rightarrow$ steepness parameter.

- ② Bipolar sigmoid fn:

Given as: $f(x) = \frac{2}{1+e^{-\lambda x}} - 1 = \frac{2(1-e^{-\lambda x})}{1+e^{-\lambda x}} = \frac{1-e^{-\lambda x}}{1+e^{-\lambda x}}$



Sigmoid fn are so called as (z their graphs are S-shaped).

Sigmoid fn defined to be odd, are monotonic fn of z our variable.

* 3) Linear actvn fn:

Defined as, $f(x) = x$ for all x .

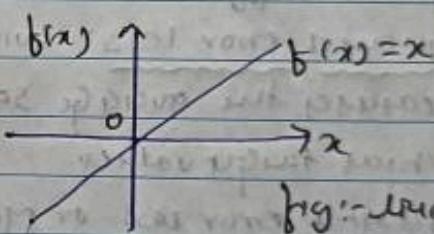


fig. Linear actvn fn.

* 4) Softmax activation fn:

- also known as softmax or normalized exponential fn. It converts a vector of K real no. into a probability distribution of K possible outcomes.

- It is a generalization of logistic fn; forwardly,

$$\therefore \sigma = \text{softmax} ; \vec{z} = \text{i/p vector} \quad \sigma = (\vec{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

e^{z_i} = stand. exponential fn for i/p vector.

K = no. of classes in multi-class classification.

e^{z_i} = standard exponential fn for o/p vector.

5) ReLU fn: Rectified linear unit.



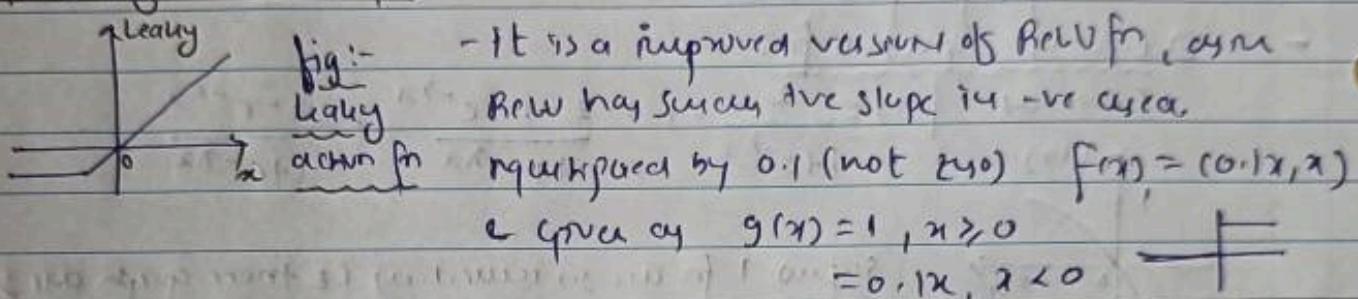
Don't activate all the neurons at the same time and vanishing gradient problem doesn't work in ReLU.

ReLU: $f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$

fig:- ReLU
actn fn.

$\therefore f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$

* 6) Leaky ReLU fn:-



5. Explain different loss fn in detail.

- > 1) Squared error loss: widely used loss fn in ML and statistic.
 - measures the average squared difference b/w predicted values & the actual target values.
 - Squared error loss or Mean Squared error loss is calculated by squaring the difference b/w the value by & the predicted value, we sum the no. to get a total value & then divide the no. by again.
 - Thus, $MSE = \frac{1}{n} \sum (y - \hat{y})^2$... this yields a no.

> 2) Cross-Entropy

- Binary cross entropy: This loss fn is used for binary classification tasks, it measures the dissimilarity b/w predicted probability & true binary labels.

Given as:- $\frac{1}{N} \sum_{i=1}^N (y_i \log(\hat{y}_i) + (1-y_i) \log(1-\hat{y}_i))$ \Rightarrow Binary class.

- Categorical cross entropy: for multiple-class classifn & categorical cross entropy is used. It measures the dissimilarity b/w predicted class



probabilities and one-hot encoded true labels.

$$\text{Spred loss} : - \sum_{i=1}^n y_i \log(g_i) \rightarrow \text{for multi-class.}$$

> 3) choosing opf fn and loss fn:

- ① when Numerical value of opf, which you want to predict.
 - activfn: binary activation fn. & ReLU
 - loss fn: MAE
- ② when the opf you trying to predict is binary.
 - activfn: sigmoid fn
 - loss fn: Binary cross entropy fn.
- ③ when the opf you trying to predict for multiple class.
 - activfn: softmax fn
 - loss fn: categorical cross entropy loss fn.

6. Explaining Gradient Descent, stochastic GD and minibatch GD in detail.

Gradient Descent is an optimization algorithm and it is used to train ML models and Neural Netw.

Training the data helps these models learn over time and the cost fn upto which gradient descent reaches its accuracy, with each step of parameter update.

- The working of gradient descent is to reduce the cost fn or the error between predicted & actual value.

- To do this, we require two data points: a direction & a learning rate.

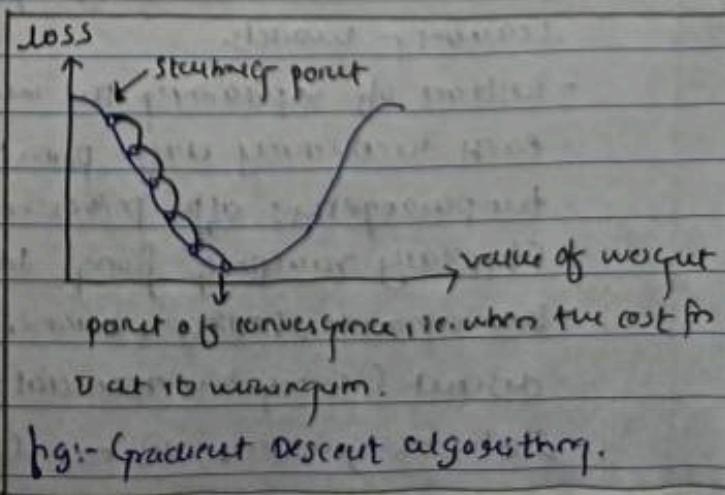


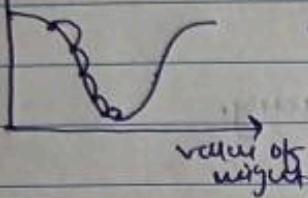
fig:- Gradient descent algorithm.



Learning rate: also called as step size or alpha. This is the size of the steps towards minimum.

① Large learning rates: Result in large steps but the risk involved is running its potency depends on the behaviour of loss fn.

② \downarrow loss



a large learning rate may

small step size, advantage

of loss function is give more precision. since we.

fig:- small learning rate. of stream are more...

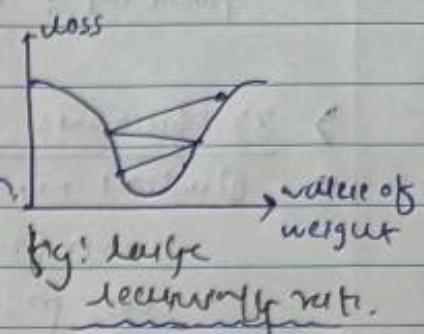


fig: large learning rate.

* Stochastic Gradient Descent :- SGD is an optimization algorithm commonly used in training ML models, including deep neural networks.
- Unlike traditional GD which calculates the average gradient over the entire dataset before updating the model's parameters, SGD updates the parameters after processing each individual data point (or a small subset called minibatch). This introduces randomness into updates, which help escape local minima & make training process faster.

* mini-batch Gradient Descent :- mini-batch GD is a variation of the stochastic gradient descent (SGD) optimization algorithm used in training deep learning models.

- Instead of updating the model's parameters after processing each individual data point (as in SGD), mini-batch GD updates the parameters after processing a small batch of data points (typically ranging from tens to hundreds).

- This approach strikes a balance between the efficiency of batch gradient descent (using the full dataset) and the noise introduced by the stochastic gradient descent (SGD).

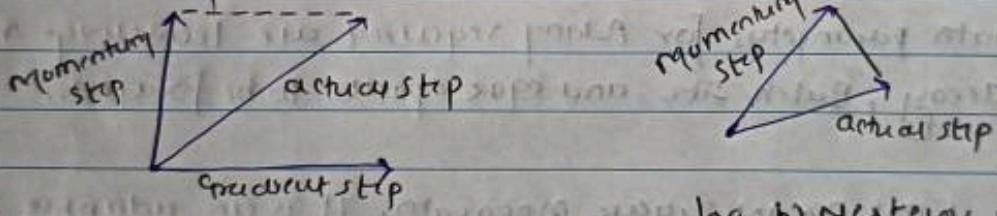


7. Write a short note on :

- * Momentum based GD - momentum based Gradient descent is an optimization algorithm used to train different models.
- The learning process consists of an objective fn (or error fn), which finds out error a ML model has on given dataset.
- The momentum term is set to a value b/w 0 and 1 with a higher value resulting in a more stable optimization process.
- Advantage: momentum is faster than SGD and the learning will be faster than SGD.

* Nesterov Accelerated GD - NAG is a modified version of momentum with stronger theoretical convergence and guarantees for convex fn.

- Momentum vs. Nesterov



(a) Momentum update

(b) Nesterov momentum update

- In the standard momentum method, the gradient is computed using current parameters (θ_t).
- Nesterov momentum achieves stronger convergence by applying the velocity (v_t) to the parameters in order to compute intertemp parameters ($\theta = \theta_t + \mu v_t$), where μ is the decay rate.
- Advantage of NAG is that it allows larger decay rate.

* Adagrad : 'Adagrad' stands for 'Adaptive Gradient Optimizer'. The optimizers like gradient descent, stochastic gd, mini-batch SGD are used to reduce the loss fn with respect to the weights.



→ The weights update formula is given by,

$$(w)_{\text{new}} = w(\text{old}) - \eta \frac{\partial L}{\partial w(\text{old})}$$

... η = learning rate.
 $w(t)$ → value of w at current iteraⁿ

for iteraⁿ, the formula, $w_t = w_{t-1} - \eta \frac{\partial L}{\partial w(t-1)}$... $w(t-1)$ at previous iteraⁿ.

-advantages : fast convergence & more stable.

* Adam: The name 'Adam' is derived from 'adaptive momentum estimation'. This optimization algorithm is an extension of SGD to update new weights during training.

- Benefits of using Adam on non-convex optimization problems are:
 - > straight forward to implement.
 - & > computationally efficient.
- Adam is known for its fast convergence & ability to work well on noisy and sparse datasets.
- Configuration parameters for Adam required are learning rate, weight decay, Batch size and max epochs. to be fixed..

* RMS Prop: Root mean square propagation. It is an adaptive learning rate optimization algorithm. Extension of popular Adaptive Gradient algorithm and is designed to adaptively reduce the amount of computational effort used in learning neural networks.

- In RMS Prop, each update is done according to can described below. This update is done separately for each parameter.

for each parameter w_j (we drop superscript j for clarity)

$$v_t = p v_{t-1} + (1-p) g_t^2$$

$$\Delta w_t = - \frac{\eta}{\sqrt{v_t + \epsilon}} g_t \quad \dots \eta \rightarrow \text{fixed learning rate.}$$

$$w_{t+1} = w_t + \Delta w_t \quad \dots \text{eqn 3}$$

$v_t \rightarrow$ Exponential avg of square of gradients.

$g_t \rightarrow$ gradient at time t along w_t .

8.

Explain the terms:

- * Parameter sharing: (weight sharing) \rightarrow many set of parameters to be shared as we interpret various models or model components as sharing a unique set of parameters. We only store a subset of memory. Eg: Natural language having specific statistical properties that are robust for translation.
- * Weight decay: weight decay is a regularizer technique that is used in ML to reduce the complexity of a model & prevent overfitting.
 The new loss function with weight decay technique is:

$$L(w, b) + \lambda/2 \|w\|^2.$$
 here $L(w, b)$ represents the original loss for before adding regularization L_2 , weight (weight decay) term.
- * Dropout: The typical characteristic of CNN is dropout layers. The dropout layer is a mask that nullifies the contribution of some neurons towards the next layer and leaves unmodified all others.
- * Early stopping: In ML, early stopping is a form of regularization used to avoid overfitting while training a model with an iterative method, such as gradient descent.
- * Data Augmentation: Data augmentation is a technique of artificially increasing the training set by creating modified copies of a dataset using existing data.
 Techniques: 1) Audio data augmentation
 2) Text data augmentation
 3) Image augmentation...

★ Bias-variance tradeoff: The bias-variance tradeoff is a fundamental concept in ML that refers to the balance b/w two sources of error that affect the performance of a predictive model: bias & variance. Finding the right balance b/w these two sources of error is crucial for building a model that generalizes well to unseen data.

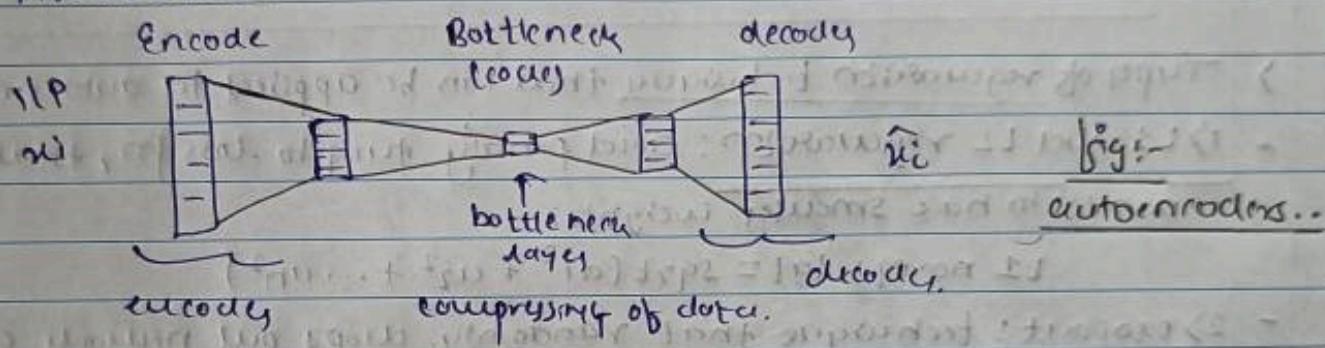
Q. What is Regularization?

- > Regularization is a set of techniques used in ML to prevent overfitting, which occurs when a model learns to fit the training data too closely, capturing noise and making it perform poorly on unseen data.
- > How does regularization help in reducing overfitting?
 - 1) penalizing complexity: Regularization methods penalize complex models by adding a cost to the loss function based on the complexity of model. This discourages the learning model from fitting noise in the data.
 - 2) constraining parameters values: Some regularization techniques, like weight decay (L2 regularization), restrict the magnitude of parameter values, preventing them from becoming too large. This encourages the model to have smaller, more manageable weights.
 - 3) encouraging simplicity: Techniques like dropout and early stopping force the model to learn more robust and simple representation by randomly dropping neurons or stopping learning when performance on validation data degrades.
- > By imposing these constraints on penalties, regularization techniques bias the model toward simpler solutions that generalize better to unseen data, effectively reducing overfitting.

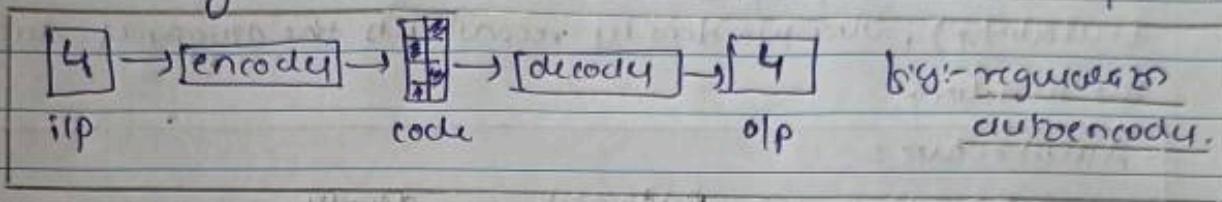
10. Explain autoencoders : (Autoencoders)

- > Autoencoders are neural nets used for unsupervised learning and dimensionality reduction. They consist of an encoder and a decoder. The encoder maps input data to a lower-dimensional representation (encoding), and the decoder reconstructs the original data from the encoding.

- Architecture:



- > Justify the advantages of autoencoders over principle component analysis for dimensionality reduction.
- ① Non-linearity: Autoencoders can capture complex, non-linear relationships in the data, while PCA is inherently linear.
- ② Representation learning: Autoencoders learn meaningful features directly from the data, whereas PCA relies on orthogonal transformation.
- ③ Flexibility: Autoencoders can handle various types of data, including images, text and sequences, while PCA is primarily designed for numerical data.
- ④ Robustness: Autoencoders can handle noisy data & perform well when data has missing values, which can be challenging for PCA.
- > Autoencoders are typically more computationally intensive & require large amounts of data for training compared to PLA. The choice between methods depends upon specific problem and data characteristics.

11. Explain regularization in autoencoders & its techniques in detail.
- > Regularization in autoencoders is used to prevent overfitting, just like in other machine learning models. Overfitting in autoencoders can lead to reconstructs that are overly robust on noise or simply redundant in training data.
- 

```

graph LR
    I[Input IIP] --> E[encoder]
    E --> C[code]
    C --> D[decoder]
    D --> O[Output OIP]
    style C fill:#ccc,stroke:#000,stroke-width:1px
    style D fill:#ccc,stroke:#000,stroke-width:1px
    style E fill:#ccc,stroke:#000,stroke-width:1px
    style O fill:#ccc,stroke:#000,stroke-width:1px
    C -- "big-regression" --> D
    C -- "autoencoder" --> O
    
```
- > Types of regularization techniques that can be applied to autoencoder:
- 1) L1 and L2 regularization: adds penalty term to loss fn, that encourages model to have smaller weights.
 $L1 \text{ norm: } \|a\|_1 = \sqrt{a_1^2 + a_2^2 + \dots + a_n^2}$
 - 2) Dropout: technique that randomly drops out neurons during training to prevent overfitting.
 - 3) Batch normalization: Technique that normalizes the IIP to each layer to have zero mean & unit variance. This can help to improve stability of training process & prevent overfitting.
 - 4) Data augmentation: Involves applying transformation to IIP data, such as cropping or flipping, to increase size of training set and improve generalization performance. Eg: for images.
 - 5) Geometric transformations: You can randomly flip, crop, rotate or translate image, & this is just the tip of the iceberg.
 - 6) Color space transforms: change RGB color channels, intensify any colors.
 - 7) Kernel filters: sharpen or blur an image.
- > Regularization techniques can be effective way to improve performance of autoencoders & prevent overfitting. The choice of regularization technique will depend on specific problem & data set being considered.



Name :- Prathamesh 5. Chikankar

ROLL NO. :- AIML11

BRANCH :- (SE-(AI&ML))

Year:- BE Subject :- Deep learning (DL)

Topic:- Assignment No. 02

SIGN:- (Pratna)

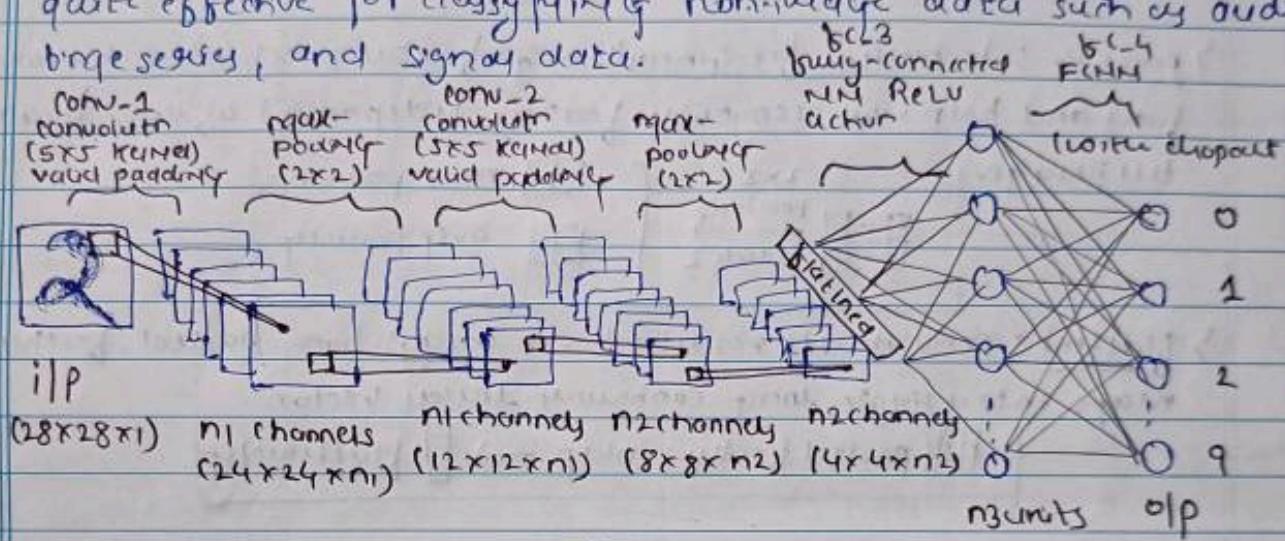
Date:- October '23

Pratna
Oct 23

A

1. Explain CNN architecture in detail.

A Convolutional Neural Network (CNN), is a neural architecture for DL which learns directly from data. CNN are particularly used for finding patterns in images to recognize objects. They can also be quite effective for classifying non-image data such as audio, time series, and signal data.



Key: CNN architecture

Kernel or Filter or Feature Detectors: filter that is used to extract the features from the images.

$$\text{formula} = [i-k] + 1 \quad \dots \quad k \rightarrow \text{size of kernel}$$

$i|p \rightarrow$ image patch at kernel $\rightarrow o|p$
(local receptive field) (filter)

Stride: modify the amount of movement over the image or video.

If stride is 1 then 1 by 1. If stride is 2 it skips next 2 pixels.

$$\text{formula} = [i-k/s] + 1 \quad \dots \quad s \rightarrow \text{stride}$$

$7 \times 7 i|p \text{ volume} \rightarrow$ $5 \times 5 o|p \text{ volume}$
stride 2

$7 \times 7 i|p \text{ volume} \rightarrow$ $3 \times 3 o|p \text{ volume}$
stride 2

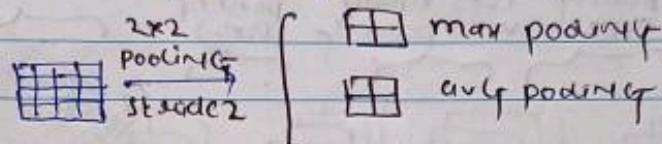


padding: being relevant to CNN as it refers to the no. of pixels added to an image when it is being processed by kernel of CNN.

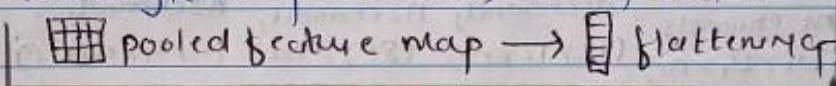
$$\text{formula} = [(w - k + 2p)/s] + 1 \quad \dots p \rightarrow \text{padding}$$

$$\text{image} * [\text{filter} * \text{kernel}] = \text{feature}$$

pooling: technique for generalizing features extracted by convolutional filters and help NLP recognize feature independent of their location in the image.



flattening: convert all resultant 2-D arrays from pooled feature maps into a single long continuous array vector.



layers used to build CNN:

- convolutional layers

- fully-connected layers

- pooling layers

- (FC)

- convolution layers

- pooling layers

$$w - f + 2p/s + 1$$

$$w - F/s + 1$$

dropout: Nullifies contribution of some neurons towards next layer and leaves unmodified all others.

activation fn: decide whether a neuron should be activated or not.

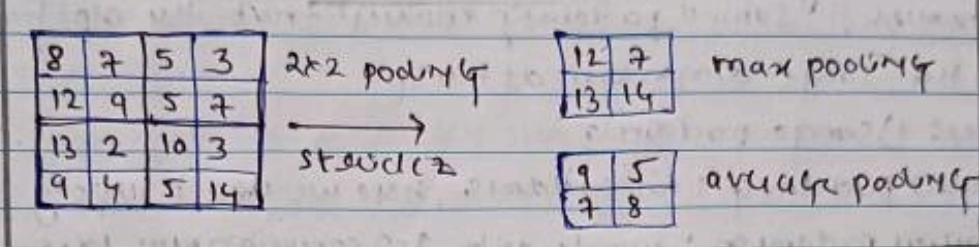
This means that it will decide whether the neuron is important to NLP or NLP is important or not in process of prediction.

They are severely commonly used activation fn such as ReLU, softmax, tanH, and the sigmoid fn.

2.

Explain the terms w.r.t. CNN with example.

> Pooling: pooling in CNN is a technique for generalizing features extracted by convolutional filters and helping the NN to recognize features independent of their location in the image.



Example: selects the maximum value from a group of numbers, if we apply 2x2 max pooling to a 4x4 feature map, output size becomes 2x2.

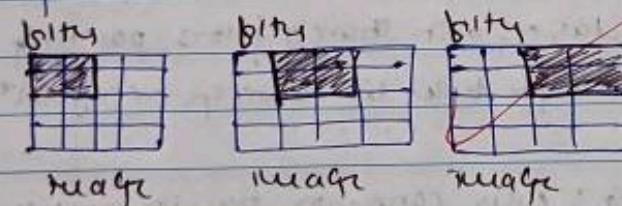
> Stride: stride is a parameter of NN filters for modifying elements of receptive field over the image or video. We had stride 1 so it will take one by one. If we give stride 2 then it will take value by skipping the next 2 pixels.

$$\text{formula} = [i - k(s) + 1]$$

$i \rightarrow$ index of ip

$k \rightarrow$ size of kernel

$s \rightarrow$ stride



left image: stride 0

middle image: stride 1

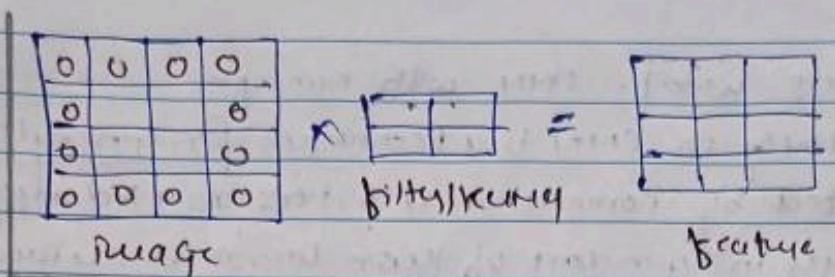
right image: stride 2

↳:- convolution (stride)

size of op maps calculated by,

$$\left[\frac{(n+2p-b+1)}{s} + 1 \right] \times \left[\frac{(m+2p-b+1)}{s} + 1 \right]$$

> Padding: padding is a term relevant to CNN as it refers to the no. of pixels added to an image when it is being processed by the ^{kernel} of a CNN.



for example, "same" padding ensures that the output feature map has the same dimensions as the ip.

Types: 1) same padding

2) valid padding: no padding, here we don't apply any padding

3) causal padding: works w/ 1-D convolutional layers.

3. Explain how fully connected NN is different from CNN.

1) Data type:

FCN: for structured data, like spreadsheets or numerical data.

CNN: for grid-like data, especially images & spatially correlated info.

2) layer types:

FCN: Each neuron connects to all neurons in next layer. No shared weights.

CNN: uses convolutional layers with shared filters, pooling layers, and fully connected layers for tasks like image recognitn.

3) weight sharing:

FCN: no weight sharing; every connection has its own weight.

CNN: uses weight sharing, allowing same filter to be applied at different posns in ip, which is robust to shifts.

4) spatial hierarchy:

FCN: doesn't capture spatial hierarchy or local patterns.

CNN: specially designed to detect local patterns & hierachical features in grid-like data.

5) translation invariance:

FCN: lacks inherent translation invariance, treats smtch features in

different locate differently.

RNN: inherently possess translation invariance, making it suitable for recognizing patterns regardless of their location in an image.

6) USE CASES:

RNN: ideal GPU for structured data analysis & non-sparsity tasks like text classifiers.

CNN: ideal for image-related tasks, particularly image classification, object detection & image segmentation.

4. Difference between

Aspect	LeNet	AlexNet
year of introduction	1998	2012
architecture depth	relatively shallow for its time	considered deep at the time of release
convolution layers	2	5 (including some larger filters)
pooling layers	2 (Avg pooling)	3 (max pooling)
activation fn	Sigmoid	Rectified linear unit (ReLU)
local response normalization (LRN)	not used	used in specific layers to reduce overfitting
dropout	not used	used to prevent overfitting.
fully connected layers	3 (including the output layer)	2 (including the output layer)
size of filters	smaller (eg., 5x5 & 3x3)	larger (eg., 11x11, 5x5, 3x3)
no. of filters	fewer (eg., 6 and 16)	more (eg., 96, 256 & 384)
use of GPU	initial GPU use (less common)	utilized GPUs extensively
winner of ImageNet	No	Yes (won the ImageNet competition).

5. Explain the architecture of ResNet in details with its advantages.
- repeating Residual blocks are stacked together to form a ResNet. we have "skipped connections" which are the major part of ResNet. The following figure provided denotes how residual block works. The idea is to connect the ip of a layer directly to the op layer after skipping & a few connector.

$[7 \times 7, 64, s=2] \quad p=3$

$[3 \times 3 \text{ maxpool}] \quad s=2$

$\left\{ \begin{array}{l} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{array} \right\} x_3$

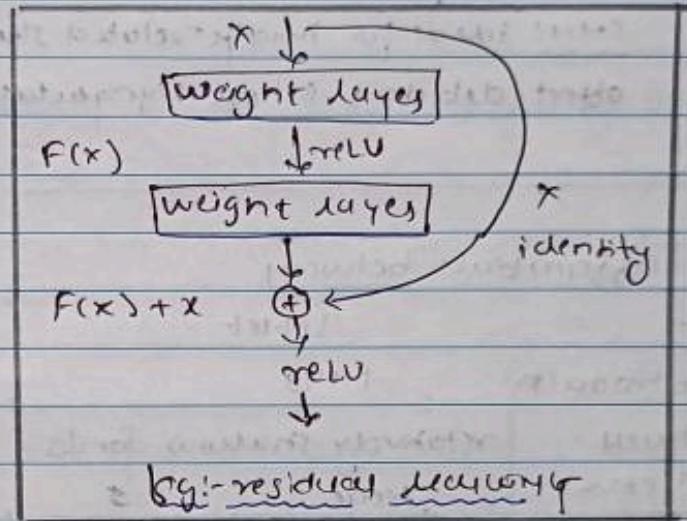
$\left\{ \begin{array}{l} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{array} \right\} x_4$

$\left\{ \begin{array}{l} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{array} \right\} x_5$

$\left\{ \begin{array}{l} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{array} \right\} x_6$

$\boxed{\text{Avg pooling, } 1000 = f(x)}$

~~30% max~~



$$\alpha^{[1]} = x$$

$$\alpha^{[l+1]} = w^{[l+1]} \alpha^{[l]} + b^{[l+1]}$$

$$\alpha^{[l+1]} = \text{relu}(\alpha^{[l+1]})$$

$$\alpha^{[l+2]} = w^{[l+2]} \alpha^{[l+1]} + b^{[l+2]}$$

$$\alpha^{[l+2]} = \text{relu}(w^{[l+2]} \alpha^{[l+1]} + b^{[l+2]} + \alpha^{[l]})$$

$$= \text{relu}(\alpha^{[l+2]} + \alpha^{[l]})$$

We can see here x is the ip to the layer which we are directly using to connect to layers after skipping the identity connection and if we think the op from identity connection to be $f(x)$. Then we can say the op will be $F(x) + x$.

- input: $n \times n$
- padding: p
- filter size: $f \times f$
- output: $((n+2p-f)/s+1) \times ((n+2p-f)/s+1)$

advantages:

- each block augments the data (additional data for learning)
- facilitate training of very deep NN w/o vanishing gradient problem.
- shorty gradient path (inception)
- active state of art performance on various NLP tasks.
- modularity
- reduces the need of extremely large datasets.

8. Differences b/w

Aspect	Recurrent Neural Network (RNN)	Feed Forward Neural Network (FNN)
Architecture	Recurrent connect, allowing information to circulate through the HNs over time.	No recurrent connect; information flows one way from IP to OP.
Data type	Designed for sequential data, such as image series, text, or speech	Typically used for structured data like tabular form, numerical data, & non-sequential data.
Layer type	Consist of recurrent layers, such as Simple RNN, LSTM, or GRU	composed of IP hidden & OP layers, with no recurrent connect.
Hidden state	Repeating hidden state that stores info. from previous time steps.	No hidden state; each IP is processed independently.
Temporal processing	Handles sequences by considering the order and temporal dependency in data.	Does not inherently capture temporal dependency.



Output length flexibility	Can handle i/p sequence of varying lengths & produce o/p at each time step	Typically produces a fixed-length o/p for a given i/p.
Use case	Time series predictor, NLP, speech recognition, etc..	Classification, regression and structured data analysis where sequence i/p problem is not required.
Vanishing gradient	Prone to vanishing gradient problems when processing long sequences.	Less prone to vanishing gradient issues.
Training method	Trained using backpropagation through time (BPTT) for sequence learning.	Training using standard backpropagation.
Example appn	Language translation (eg: with seq2seq models), text generation, sentiment analysis	Fimage classification, recommendation systems, financial modeling.

9.

Explaining RNN architecture in detail.

A recurrent neural net is a type of deep learning neural net that remembers the i/p sequence, stores it in memory states/cell states, and predicts the future word/sequence sentences.

Let's consider x_{11}, x_{12}, x_{13} as i/p & o_1, o_2, o_3 as o/p of hidden layers 1, 2 and 3 respectively. The i/p are sent to the RNN at different time intervals, so lets say x_{11} is sent to the hidden layer 1 at time t_1 , x_{12} at t_2 , x_{13} at t_3 .

Also, let's assume weights are same in forward propagation.

The o/p o_3 is dependent on o_2 which in turn is dependent on o_1 .

$o_1 = f(x_{11} * w)$ → where w is weight and f is activation fn.

$$o_2 = F(o_1 + x_{12} * w); \quad o_3 = f(o_2 + x_{13} * w)$$

Finally, the o/p of O_3 is the actual o/p predicted by \hat{y} .

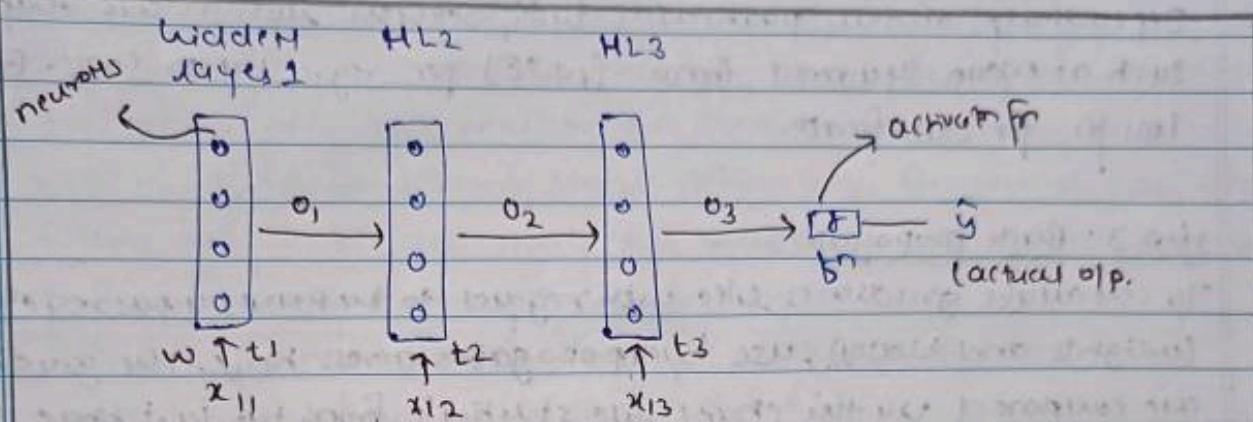


Fig:- architecture & working of simple RNN

The predicted value \hat{y} is nothing but O_3

$$\text{i.e. } O_3 = f((x_{11} \times w) + (x_{12} \times w) + (x_{13} \times w))$$

Applications: speech recognition (Google voice search)

machine translation (Google Translate)

time series forecasting

sales forecasting, etc..

10. Explain Backpropagation through time (BPTT) algorithm in detail.

Step 1: forward pass

for the whole o/p sequence, unroll the RNN over time, and calculate the forward pass for each time step. The hidden state is updated throughout the forward pass, and the RNN is equipped to compute the o/p at each time step.

Step 2: calculating the loss

By contrasting the expected o/p (y_t) with the actual target (the ground truth) at each time step, determine the loss at each



time step. The sum of the losses at all time steps is the overall loss. Depending on the particular task, several loss functions are employed such as mean squared error (MSE) for regression or cross-entropy loss fn. for classification.

Step 3: Back propagation:

To calculate gradients w.r.t. regard to the RNN's parameters (weights and biases), use backpropagation across time. The gradients are computed via the chain rule starting from the last time step ($t=1$) and going back to the initial time step ($t=1$). As they share weights, the gradients are accumulated over time steps.

Step 4: Gradient update

using the obtained gradient, modify the RNN's parameters. For this update, you can use the common gradient descent or one of its variants (such as Adam or RMSprop). Based on the particular problem and design, an optimization technique & learning rate can be selected.

~~Challenges with BPTT.~~

- vanishing/exploding gradient problem.
- memory requirements: Greatly memory requirement.

Solutions to BPTT challenges:

- long short term memory (LSTM)
- Gated Recurrent Unit (GRU)

These topologies employ gating methods to regulate the info-flow encoding of RNN, to keep track of (such) data across longer sequences.



11. Explain vanishing gradient problem in RNN.

> As a gradient moves backward in time, its magnitude gets less and smaller until it vanishes entirely. This occurs when the recurrent weights are nearly zero or when the activator fn, such as the sigmoid fn utilized in the RNN has a small derivative. Because of this, the model has trouble updating the weights in earlier time steps, which prevents it from discovering long-term dependencies.

Consider the RNN update eqn for a single time step t.

$$h_t = f(w_h h_{t-1} + w_x x_t + b)$$

... h_t : Hidden state at time step t.

f : Activator fn (e.g. tanh or sigmoid)

w_h : weight matrix for hidden state; w_x : for ip

b : bias vector; x_t : ip at time step t.

> during backpropagation, the gradient, of the loss w.r.t. hidden state ($\partial L / \partial h_t$) is calculated & propagated backward to update the weights (w_h, w_x) and bias (b). The gradient update involves taking the derivative of the activator fn (f) and is multiplied with the chain rule as gradients are propagated through time.

> If the activator fn is the sigmoid fn ($f(z) = 1 / (1 + \exp(-z))$), its derivative is given by, $f'(z) = f(z) * (1 - f(z))$. As the absolute value of $f(z)$ approaches two, the gradient vanishes, leading to vanishing gradient problem.

* Exploding- gradients in RNN.

> On the other hand, explosive gradients happen when a gradients magnitude increases exponentially during backpropagation. This may occur if the activator function has a large derivative or if recurrent weights are substantial. The model diverges during training as a



Result of the steep gradients heavy weight updates.

- > The recurrent weights ($w-h$) at each time step, on the other hand are doubled when gradients are propagated backward in time.
- > The exploding gradient problem occurs when the gradients expand exponentially as they progress backward through sequence when the recurrent weights are significant (more than 1).

12. Explain LSTM in detail.

- > Long short-term memory, sometimes known by LSTM, is a sort of sophisticated Recurrent Neural Network (RNN) architecture. The disappearing and exploding gradient difficulties, as well as the challenge of capturing long-term relationships in sequential data are some of the shortcomings of conventional RNNs that LSTM were created to overcome.
- > Simply put, the LSTM introduces a unique memory cell and three gating mechanisms (inp gate, forget gate, and o/p gate) that allows it to selectively keep or forget info. over time. As a result, LSTM can handle long-range relationship and keep crucial info. over lengthy sequences.
 - memory cell: In LSTM many building blocks is a memory cell. It is updated using i/p, forget & O/p gates and store the content or data from every time steps.
 - i/p Gate: The gate decides which data from the most recent concealed state and the current time step should be added to the memory cell.
 - forget Gate: The forget gate makes the decision as to which data in memory cell should be deleted or forgotten.
 - O/p Gate: The o/p gate regulates how much of the data stored in the

memory sheet cell should be used to create the o/p or hidden state for the current time step.

* GRU: Gated Recurrent Unit

A modification of the conventional long short-term memory (LSTM) architecture known as the Gated Recurrent Unit (GRU) was created in order to solve some of the LSTM's complications while preserving its capacity to identify long-term dependencies in sequential data.

> Key components of GRU

- Update Gate (z_t): The update gate distinguishes how much of the previous hidden state to retain ($1 - \text{forget}$) and how much of the new candidate activation to consider (update)
- Reset Gate (r_t): The reset gate distinguishes how much of previous hidden state should be forgotten (reset) to account for new i/p.
- Candidate activation ($h_{\text{candidate}}$): The candidate activation represents the new info. that can be added to hidden state.

> Advantages of Gated Recurrent Unit (GRU) over LSTM

- simple architecture: fewer gating mechanisms than LSTM, making it easier to use & more efficient computationally.
- faster training: can frequently be trained more quickly.
- less prone to overfitting
- simpler performance to LSTM.

13. Explain GAN architecture in detail.

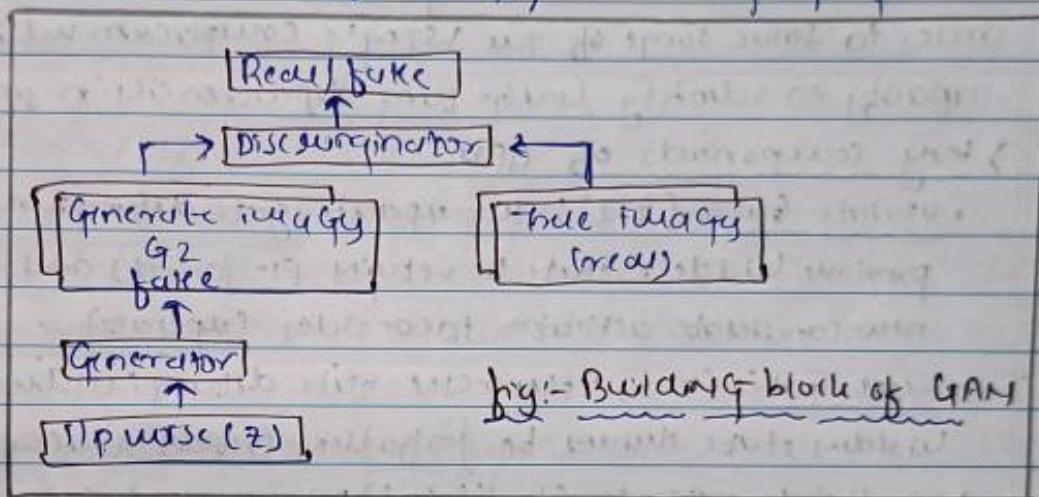
Generative Adversarial Network don't work with any explicit density estimation, but it is based on game theory approach with an objective to find Nash equilibrium, b/w two NN, generator & discarding monitor. Monitor is to sample from a simple distribution like Gaussian & compared to

turns learn to transforming this noise to data distribution using universal function approximators such as Neural Net.

> Generative: To learn generative model, which describes how data is generated in terms of a probabilistic model.

Adversarial: learning of model is done in an adversarial setting.

MIW: use DNN as AI algorithm for learning purpose.



> The objective in GAN is defined as:

$$\text{min. } \max V(D, G) = E_x \sim p(\text{data}(x)) [\log D(x)] + E_z \sim p_z(z) [\log (1 - D(G(z)))]$$

... where G = Generator ; x = Sample from $p(\text{data}(x))$

D = Dsriminator ; z = sample from $p(z)$

$p(\text{data})(x)$ = distribution of real data

$p(z)$ = distribution of generator ; $g(z)$ = generator MIW.

$D(x)$ = discriminator MIW

Approach:

• image generation: involving the use of two-part MIW - a generator and a discriminator. Generator creates fakes images while discriminator evaluates.

• deep fake: used to create realistic but fake videos or images by swapping faces or modifying content in an existing video.

6.

Given an input image size $28 \times 28 \times 1$. Apply 1 convolution layers.
 first layer - apply 64 filters each of size of 3×3 , max pooling $\downarrow - s=2, f=2$
 calculate the output volume, activation shape & no. ofs parameters for
 each layer.

Soln:-

Layer	layer Neuron shape	Activation shape	parameters.
Input \rightarrow I/p	$(28 \times 28 \times 1) - 64$ ^{input} _{image}	$28 \times 28 \times 1 = 784$	0
Conv2D (filter $= 3 \times 3$)	$(28 \times 26 \times 64)$	$43,264$	640
maxpooling 2D	$(13 \times 13 \times 64)$	$10,816$	0
flatten layer	$(10816, 1)$	(10816)	0
Dense	$(128, 1)$	128	$13,84,576$
Dense	$(10, 1)$	10	1290 (13 min)
10 neuron dense layer			

1) I/p: $28 \times 28 \times 1$ ① O/p volume = $\frac{28+0-3+1}{2} = 25+1 = 26$ //

64 filters \downarrow
 $b=3$ $\therefore (26, 26, 64) = 43,264$

② 64 filters of size $3 \times 3 \times 1$ channel $[(3 \times 3 \times 1) + 1] \times 64$
 $10 \times 64 = 640$ //

2) pooling \downarrow \rightarrow O/p vol = $\frac{26-3+1}{2} = 12.5 = 13$ //

$(26, 26, 64)$
 $s=2$

\downarrow 2×2 pooling \downarrow $\frac{26-2}{2} + 1 = \frac{24}{2} + 1 = 13$ //

Conv2D Dimensions - pool $|$ dilat $+ 1 = \frac{26-2}{2} + 1 = 13$ //

stereo

3) Dense : layer 1



(current layer memory + previous layer memory)

+ 1 + (current layer neurons)

$$(128 \times 10,816) + 1 + 128 \Rightarrow 13,84,576$$

(1.3 million)

4) Denote layer 2:

$$(10 + 128) + 1 \times 10 \Rightarrow 1290$$

7.

Given an input image volume of dimension $100 \times 100 \times 3$. For first convolutional layer, apply 8 filters of size $3 \times 3 \times 3$. For second convolutional layer apply 1 filter of size 3×3 . Calculate the no. of parameters at each convolution layer. (Assume $s=1, p=0$)

Soln:-

